



SIX WEEKS SUMMER TRAINING REPORT
ON
BIG DATA

TOPIC: Crime Recorded (2001-2014)

Submitted by

Name: Vinayak Vedant Srivastava

Course: *B.Tech (Computer Science Engineering)*

Under the guidance of

Ms. Gurpreet Kaur

**SHRI RAMSWAROOP MEMORIAL UNIVERSITY, UTTAR
PRADESH**

DECLARATION

I hereby declare that I have completed my six-week summer training at IBM under the guidance of Ms. Gurpreet Kaur. I hereby undertake that the project undertaken by me is the genuine work of mine.

Name of Student: Vinayak Vedant Srivastava

Date:.....

ACKNOWLEDGEMENT

We would like to express our deepest appreciation to all those who provided us the possibility to complete this report. A special gratitude we give to our 6th semester B.D.E. project supervisor, Ms. Gurpreet Kaur, whose contribution in stimulating suggestions and encouragement, helped us to coordinate our project and especially in writing this report. Furthermore, we could also like to acknowledge with much appreciation her crucial role, in giving the permission to use all required equipment and the necessary materials to complete the task 'Analysis of AMCAT data' using Apache Hadoop (Pig and Hive), IBM Cognos insight and R Programming, and gave suggestions about the task.

- **Vinayak Vedant Srivastava**

TABLE OF CONTENTS

1. Introduction

1.1 What is Big Data?

1.2 Types of Big Data

1.2.1 Structured

1.2.2 Unstructured

1.2.3 Semi-structured

1.3 Characteristics of Big Data

2. What is Hadoop?

2.1 Hadoop Tools

2.1.1 Hadoop Distributed File System

2.1.2 Hbase

2.1.3 HIVE

2.1.4 Sqoop

2.1.5 Pig

2.1.6 Map reduce

3. Project: Crime Recorded (2001-2014)

3.1 Purpose

3.2 Project Scope

3.3 Dataset

3.3.1 Details

3.3.2 Attributes

3.3.3 Dataset (Excel View)

4. Problem Statement

4.1 Coding

4.2 Visualization

INTRODUCTION

What is Big Data?

Big Data is also data but with a huge size. Big Data is a term used to describe a collection of data that is huge in volume and yet growing exponentially with time. In short such data is so large and complex that none of the traditional data management tools are able to store it or process it efficiently.

Types of Big Data

Big Data could be found in three forms:

1. Structured
2. Unstructured
3. Semi-structured

Structured

Any data that can be stored, accessed and processed in the form of fixed format is termed as a 'structured' data.

Unstructured

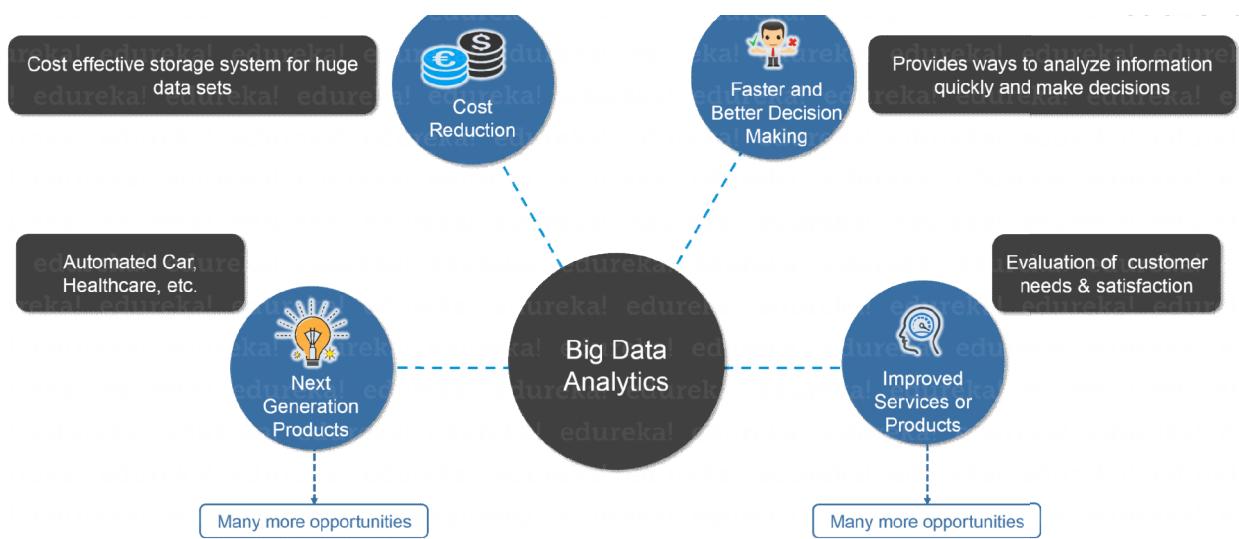
Any data with unknown form or the structure is classified as unstructured data.

Semi-structured

Semi-structured data can contain both the forms of data. We can see semi-structured data as a structured in form but it is actually not defined with e.g. a table definition in relational DBMS.

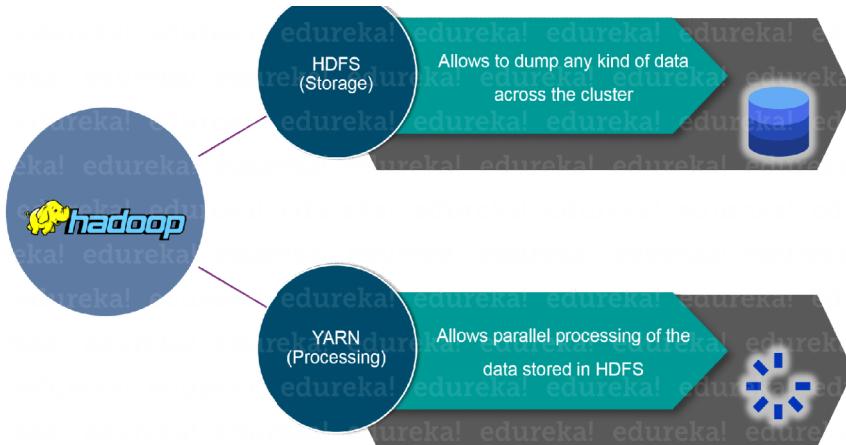
Characteristics of Big Data

- (i) **Volume**
- (ii) **Variety**
- (iii) **Velocity**
- (iv) **Variability**



WHAT IS HADOOP?

Hadoop is a framework that allows you to first store Big Data in a distributed environment, so that, you can process it parallelly.



Hadoop Tools:

1. Hadoop Distributed File System

The Hadoop Distributed File System (HDFS) is designed to store very large data sets reliably, and to stream those data sets at high bandwidth to user applications. In a large cluster, thousands of servers both host directly attached storage and execute user application tasks. By distributing storage and computation across many servers, the resource can grow with demand while remaining economical at every size. We describe the architecture of HDFS and report on experience using HDFS to manage 40 petabytes of enterprise data at Yahoo.

2. Hbase

HBase is a column-oriented database management system that runs on top of HDFS. It is well suited for sparse data sets, which are common in many big data use cases. Unlike relational database systems, HBase does not support a structured query language like SQL; in fact, HBase isn't a relational data store at all. HBase applications are written in Java much like a typical MapReduce application.

3. HIVE

The Apache Hive data warehouse software facilitates querying and managing large datasets residing in distributed storage. Hive provides a mechanism to project structure onto this data and query the data using a SQL-like language called HiveQL.

4. Sqoop

Sqoop is a tool designed to transfer data between Hadoop and relational databases. You can use Sqoop to import data from a relational database management system (RDBMS) such as MySQL or Oracle into the Hadoop Distributed File System (HDFS), transform the data in Hadoop MapReduce, and then export the data back into an RDBMS.

5. Pig

Pig is a platform for analyzing large data sets that consists of a high-level language for expressing data analysis programs, coupled with infrastructure for evaluating these programs. The

salient property of Pig programs is that their structure is amenable to substantial parallelization, which in turns enables them to handle very large data sets.

6. Map reduce

Hadoop MapReduce is a software framework for easily writing applications which process vast amounts of data (multi-terabyte data-sets) in-parallel on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.

PROJECT: RECORDED CRIME (2001-2014)

This project is to analyze a dataset related to *CRIME happened* in India from 2001-2014. This dataset contains complete information about various aspects of crimes committed in India in different regions along with the types of crime and total number of casualties.

Purpose

The purpose of this project is to analyze a dataset related to crime data along with casualties happened with the police officers. This dataset contains information about various offence committed in INDIA. It gives detailed analysis of crime and its growth in the country.

Project Scope

- The objective of this project is to analyze the rate of crimes happened in different regions of the country.
- It is valid for predictive analysis as well as descriptive analysis too.

Dataset

The dataset contains information about crimes happened in different regions of India. It also includes information about police officer killed/injured in various categories. This POC contains

Hive, Pig, Map – Reduce and Sqoop queries to analyze the data.
The dataset contains 10677 entries.

Details

This dataset contains complete information about various aspects of crimes happened in India from 2001-2014. Attributes are divided by ‘,’ i.e. they are comma separated (csv). The dataset contain:

STATE_UT,DISTRICT,YEAR,MURDER,ATTEMPT_TO_MURDER,RAPE,KIDNAPPING,KIDNAPPING_AND_ABDUCTION_OF_WOMEN_AND_GIRLS,DACOITY,ROBBERY,BURGLARY,THEFT,RIOTS,CRIMINAL_BREACH_OF_TRUST,CHEATING,COUNTERFEITING,ARSON,HURT_GREVIOUS_HURT,DOWRY_DEATHS,ASSAULT_ON_WOMEN,CRUELTY_BY_HUSBAND_OR_HIS_RELATIVES,CAUSING_DEATH_BY NEGLIGENCE,OTHER_IPC_CRIMES,TOTAL_IPC_CRIMES etc...

Attributes

- STATE_UT - Name of State and Union Territories
- DISTRICT - Name of District
- YEAR - Year in which crime was committed
- MURDER - Number of Murder Victims
- ATTEMPT_TO_MURDER - Number of Attempt to Murder Victims
- RAPE - Number of Rape Victims
- KIDNAPPING - Number of Kidnapping Victims

- KIDNAPPING_AND_ABDUCTION_OF_WOMEN_AND_GIRLS - Number of Kidnapping Victims of women and girls
- DACOITY - Number of Dacoity Victims
- ROBBERY - Number of Robbery Victims
- BURGLARY - Number of Burglary Victims
- THEFT - Number of Theft Victims
- RIOTS - Number of Riots Victims
- CRIMINAL_BREACH_OF_TRUST - Victims of Criminal Breach of Trust
- CHEATING - Victims of Cheating
- COUNTERFEITING - Number of Counterfeiting Victims
- ARSON - Number of ARSON Victims
- HURT_GREVIOUS_HURT - Victims of Hurt
- DOWRY_DEATHS - Victims of Dowry deaths
- ASSAULT_ON_WOMEN - Number of Assault happened on Women
- CRUELTY_BY_HUSBAND_OR_HIS_RELATIVES - Victims of Cruelty by husband and his family
- CAUSING_DEATH_BY NEGLIGENCE - Victims of death cause due to Negligence
- OTHER_IPC_CRIMES - Number of Victims of Other IPC Crimes
- TOTAL_IPC_CRIMES - Total Number of Crimes
- AUTO_THEFT - Crimes related to theft of Automotive Vehicles
- GENDER - Male or Female Victim

Dataset (Excel View):

Crime - Microsoft Excel

	X20	f _x	9869																				
1	STATE_UT	DISTRICT	YEAR	MURDER	ATTEMPT	RAPE	KIDNAPPI	KIDNAPP	DACOITY	ROBBERY	BURGLAR\THEFT	RIOTS	CRIMINAL CHEATING	COUNTER	ARSON	HURT_GRE	DOWRY	ASSAULT	CRUELTY	OTHER			
2	ANDHRA P	FADILABAD	2001	101	60	50	46	30	9	41	198	199	78	16	104	1	30	1131	16	149	1		
3	ANDHRA P	ANANTAPUR	2001	151	125	23	53	30	8	16	191	366	168	11	65	8	69	1543	7	118	1		
4	ANDHRA P	CHITTOOR	2001	101	57	27	59	34	4	14	237	723	156	33	209	9	38	2088	14	112	1		
5	ANDHRA P	CUDDAPAH	2001	80	53	20	25	20	1	4	98	173	164	12	37	2	23	795	17	126	1		
6	ANDHRA P	FEAST GOD	2001	82	67	23	49	26	4	25	437	1021	70	50	220	3	41	1244	12	109	2		
7	ANDHRA P	GUNTAKA	2001	3	1	0	0	0	5	2	0	162	1	0	0	3	0	1	0	1	1		
8	ANDHRA P	GUNTUR	2001	182	88	54	82	51	16	59	338	1122	244	67	300	8	43	1792	7	139	3		
9	ANDHRA P	HYDERABAD	2001	111	113	37	80	39	13	67	1155	2792	65	101	1293	24	0	3137	24	118	1		
10	ANDHRA P	KARIMNA	2001	162	85	56	67	49	27	50	218	392	220	25	243	5	33	1392	62	414	2		
11	ANDHRA P	KHAMMAI	2001	93	60	47	41	30	1	13	172	368	153	35	130	5	73	1026	17	180	1		
12	ANDHRA P	KRISHNA	2001	65	51	37	36	21	3	15	163	478	70	24	104	1	62	1985	10	208	2		
13	ANDHRA P	KURNNOOL	2001	133	72	29	47	47	6	22	155	297	84	6	126	2	5	1547	13	141	1		
14	ANDHRA P	MAHABODHI	2001	157	67	59	42	27	8	27	249	316	157	22	84	0	0	867	14	176	1		
15	ANDHRA P	MEDAK	2001	101	56	35	26	20	27	26	219	286	100	17	87	4	37	1367	26	100	1		
16	ANDHRA P	NALGONDA	2001	122	60	35	27	19	6	28	133	318	220	13	122	8	72	1132	31	188	2		
17	ANDHRA P	NELLORE	2001	89	69	46	90	80	12	16	244	608	97	20	177	3	65	1119	10	207	2		
18	ANDHRA P	INIZAMABAD	2001	106	49	21	38	21	7	22	158	234	51	61	122	1	30	1383	19	55	2		
19	ANDHRA P	PRAKASHM	2001	102	82	19	31	12	15	14	147	278	138	16	88	1	43	1266	5	140	1		
20	ANDHRA P	RANGA REDD	2001	214	95	72	106	83	24	78	1076	1296	65	67	527	3	67	2829	37	113	2		
21	ANDHRA P	SECUNDERABAD	2001	6	0	0	0	0	0	10	2	296	1	2	4	25	0	17	1	0			
22	ANDHRA P	SRIKAKULAM	2001	38	10	8	12	12	1	4	118	231	70	18	53	0	34	679	6	38	1		
23	ANDHRA P	VIJAYAWADA	2001	53	44	25	70	48	3	27	491	2057	19	34	614	10	17	1578	2	84	5		

PROBLEM STATEMENT

1. Total Murder being recorded in Uttar Pradesh during the year 2013.
2. Number of Rape victims found in various age groups.
3. Top 3 categories of Policemen killed or injured.
4. Categorization of Murder victims on the basis of Gender.
5. Year having maximum number of casualties (Policemen killed).
6. Total number of Dacoity happened in each year.
7. Top 5 states of Dowry Deaths during the year 2014.
8. Top 3 category of Auto Theft during year 2011.
9. Year wise Total Rape case reported v/s Total rape victims.
10. State wise percentage of IPC crimes recorded.

Query 1:- Total Murder being recorded in Uttar Pradesh during the year 2013.

CODING

Step i. Create a new directory with name *IBMPProject*

Step ii. Transfer data into Hadoop Distributed File System(HDFS)

- ~ In order to perform queries on HIVE we need to transfer the data to HDFS

```
[training@localhost ~]$ pwd  
[training@localhost ~]$ cd IBMPProject  
[training@localhost ~]$ hadoop fs -mkdir IBMPProject  
[training@localhost ~]$ hadoop fs -copyFromLocal  
/home/training/IBMPProject/Crime.csv /user/training/IBMPProject
```

```
File Edit View Terminal Tabs Help
[training@localhost ~]$ pwd
/home/training
[training@localhost ~]$ cd IBMProject
[training@localhost IBMProject]$ hadoop fs -mkdir IBMProject
[training@localhost IBMProject]$ hadoop fs -copyFromLocal /home/training/IBMProject/Crime.csv /user/training/IBMProject
[training@localhost IBMProject]$ hadoop fs -ls
Found 2 items
-rw-r--r-- 1 training supergroup      25 2020-07-13 04:46 /user/training/00
0000_0
drwxr-xr-x - training supergroup      0 2020-07-13 05:05 /user/training/IBMProject
[training@localhost IBMProject]$ hadoop fs -ls IBMProject
Found 1 items
-rw-r--r-- 1 training supergroup 297831 2020-07-13 05:05 /user/training/IBMProject/Crime.csv
[training@localhost IBMProject]$
```

Step iii. Open Hive shell

Step iv. Create a new database with name *crimedata*

Step v. Create a new table name *crimetatable*

- ~ Use the above database i.e. *crimedata* for creation of table.

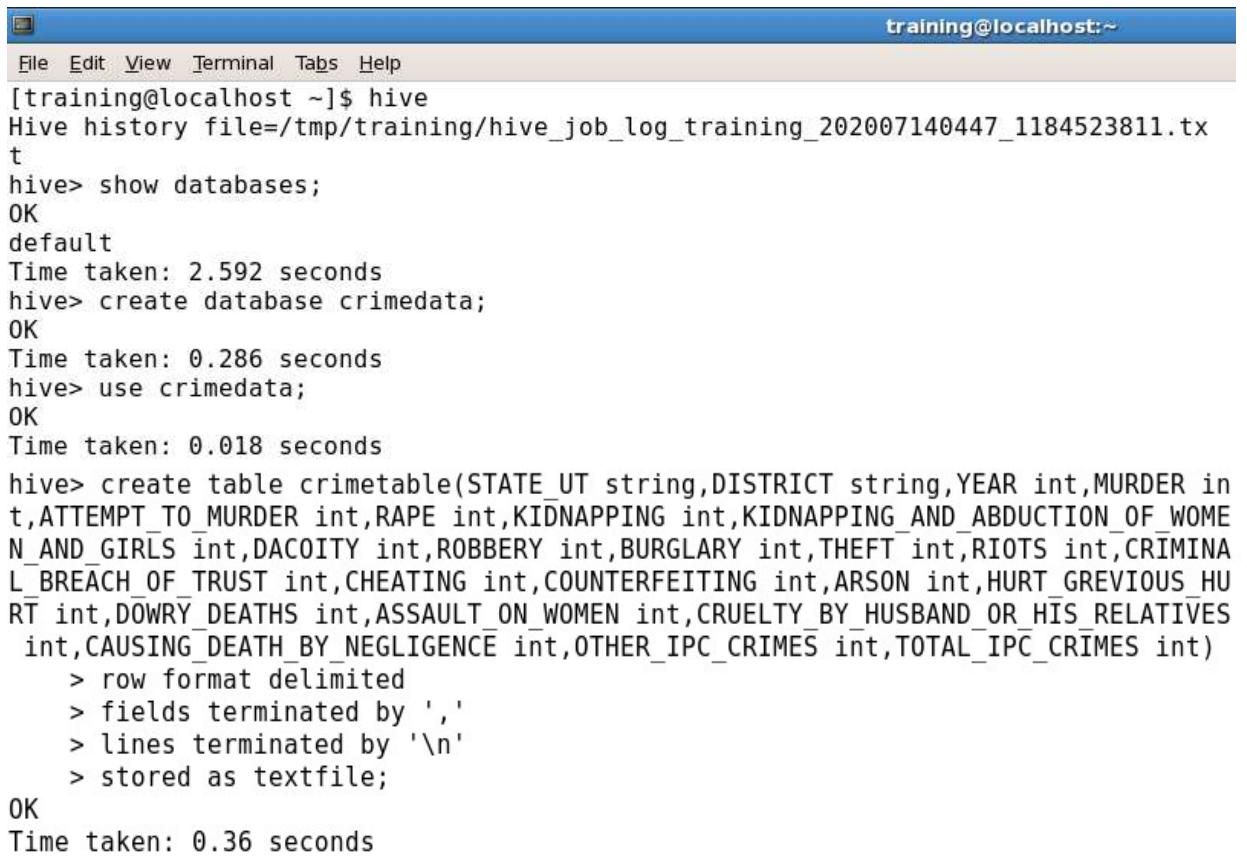
Step vi. Load the data into table *crimetatable*

```
[training@localhost ~]$ hive
hive> show databases;
hive> create database crimedata;
hive> use crimedata;
hive> create table crimetatable(STATE_UT string,DISTRICT
string,YEAR int,MURDER int,ATTEMPT_TO_MURDER
int,RAPE int,KIDNAPPING
int,KIDNAPPING_AND_ABDUCTION_OF_WOMEN_AND_
GIRLS int,DACOITY int,ROBBERY int,BURGLARY
int,THEFT int,RIOTS int,CRIMINAL_BREACH_OF_TRUST
int,CHEATING int,COUNTERFEITING int,ARSON
```

```

int,HURT_GREVIOUS_HURT int,DOWRY_DEATHS
int,ASSAULT_ON_WOMEN
int,CRUELTY_BY_HUSBAND_OR_HIS_RELATIVES
int,CAUSING_DEATH_BY NEGLIGENCE
int,OTHER_IPC_CRIMES int,TOTAL_IPC_CRIMES int)
    > row format delimited
    > fields terminated by ','
    > lines terminated by '\n'
    > stored as textfile;
hive> load data inpath '/user/training/IBMProject/Crime.csv'
into table crimetable;
hive> select * from crimetable;

```



The screenshot shows a terminal window titled "training@localhost:~". The window contains the following Hive commands:

```

File Edit View Terminal Tabs Help
[training@localhost ~]$ hive
Hive history file=/tmp/training/hive_job_log_training_202007140447_1184523811.tx
t
hive> show databases;
OK
default
Time taken: 2.592 seconds
hive> create database crimedata;
OK
Time taken: 0.286 seconds
hive> use crimedata;
OK
Time taken: 0.018 seconds
hive> create table crimetable(STATE_UT string,DISTRICT string,YEAR int,MURDER in
t,ATTEMPT_TO_MURDER int,RAPE int,KIDNAPPING int,KIDNAPPING_AND_ABDUCTION_OF_WOME
N_AND_GIRLS int,DACOITY int,ROBBERY int,BURGLARY int,THEFT int,RIOTS int,CRIMINA
L_BREACH_OF_TRUST int,CHEATING int,COUNTERFEITING int,ARSON int,HURT_GREVIOUS_HU
RT int,DOWRY_DEATHS int,ASSAULT_ON_WOMEN int,CRUELTY_BY_HUSBAND_OR_HIS_RELATIVES
int,CAUSING_DEATH_BY NEGLIGENCE int,OTHER_IPC_CRIMES int,TOTAL_IPC_CRIMES int)
    > row format delimited
    > fields terminated by ','
    > lines terminated by '\n'
    > stored as textfile;
OK
Time taken: 0.36 seconds

```

```

hive> load data inpath '/user/training/IBMPProject/Crime.csv' into table crimetable;
Loading data to table crimedata.crimetable
OK
Time taken: 0.442 seconds
hive> select * from crimetable;

```

Step vii. Again create a new table with name *murdertable*

- ~ Filtering of data is done in this phase

Step viii. Load data into murdertable by accessing data from crimedata

```

hive> create table murdertable(STATE_UT string,MURDER int)
> row format delimited
> fields terminated by ','
> lines terminated by '\n'
> stored as textfile;
hive> insert overwrite table murdertable select
STATE_UT,MURDER from crimetable where YEAR='2013';
hive> select sum(MURDER) from murdertable where
STATE_UT='UTTAR PRADESH';

```

LAKSHADWEEM	Lakshadweep	2014	0	0	1	0	0	0	0
5	1	0	1	0	30	74			
LAKSHADWEEM	Total	2014	0	0	1	0	0	0	10
1	0	1	0	30	74				
PUDUCHERRY	Karaikal	2014	6	1	3	1	0	1	1
0	8	0	12	44	322	461			
PUDUCHERRY	Puducherry	2014	19	14	7	12	1	2	5
1	97	1	20	165	1014	1890			
PUDUCHERRY	Total	2014	25	15	10	13	1	3	61
105	1	32	4	209	1336	2351			

Time taken: 0.582 seconds

```

hive> create table murdertable(STATE_UT string,MURDER int)
> row format delimited
> fields terminated by ','
> lines terminated by '\n'
> stored as textfile;

```

OK

Time taken: 0.042 seconds

```
hive> insert overwrite table murderstable select STATE_UT,MURDER from crimetype where YEAR='2013';
Total MapReduce jobs = 2
Launching Job 1 out of 2
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_202007140415_0005, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007140415_0005
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007140415_0005
2020-07-14 05:46:01,568 Stage-1 map = 0%,  reduce = 0%
2020-07-14 05:46:03,586 Stage-1 map = 100%,  reduce = 0%
Ended Job = job_202007140415_0005
Ended Job = -230666345, job is filtered out (removed at runtime).
Moving data to: hdfs://localhost/tmp/hive-training/hive_2020-07-14_05-45-59_133_6982992251250274676/-ext-10000
Loading data to table crimedata.murderstable
Deleted hdfs://localhost/user/hive/warehouse/crimedata.db/murderstable
Table crimedata.murderstable stats: [num_partitions: 0, num_files: 1, num_rows: 0, total_size: 11338]
822 Rows loaded to murderstable
OK
Time taken: 5.864 seconds

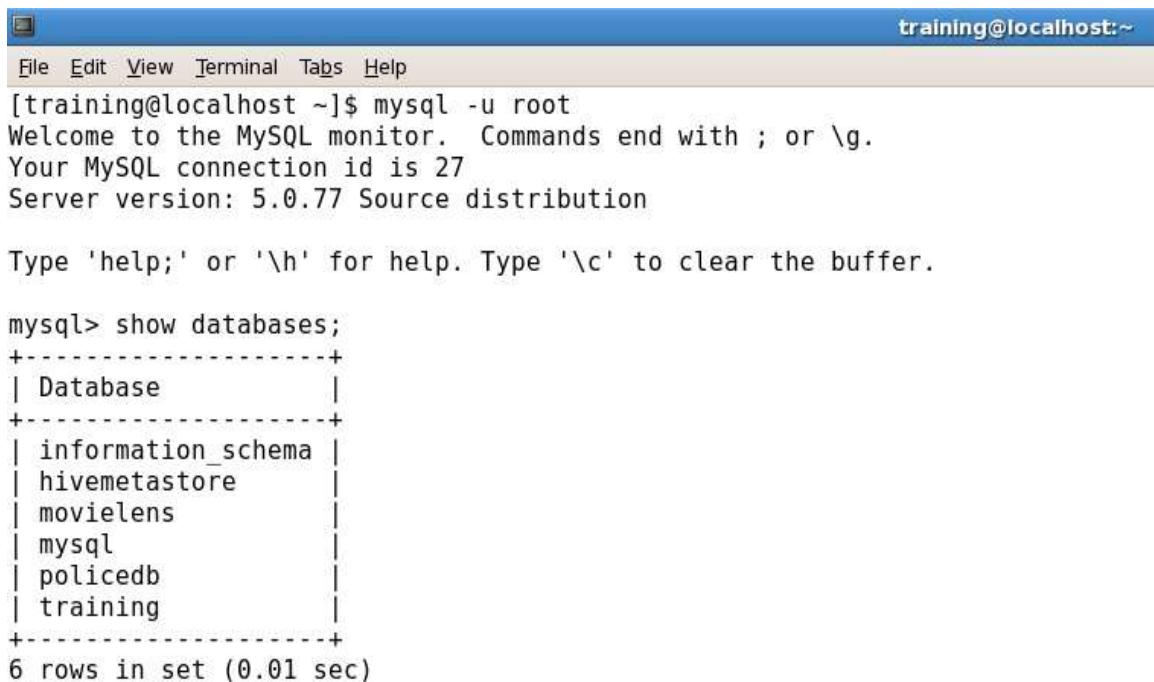
hive> select sum(MURDER) from murderstable where STATE_UT='UTTAR PRADESH';
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202007140415_0006, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007140415_0006
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007140415_0006
2020-07-14 05:47:43,900 Stage-1 map = 0%,  reduce = 0%
2020-07-14 05:47:46,014 Stage-1 map = 100%,  reduce = 0%
2020-07-14 05:47:54,343 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_202007140415_0006
OK
10094
Time taken: 14.989 seconds
```

Query 2:- Number of Rape victims found in various age groups.

CODING

Step i. Open mysql

```
[training@localhost ~]$ mysql -u root  
mysql> show databases;
```



A screenshot of a terminal window titled "training@localhost:~". The window shows the MySQL monitor with the following output:

```
File Edit View Terminal Tabs Help  
[training@localhost ~]$ mysql -u root  
Welcome to the MySQL monitor. Commands end with ; or \g.  
Your MySQL connection id is 27  
Server version: 5.0.77 Source distribution  
  
Type 'help;' or '\h' for help. Type '\c' to clear the buffer.  
  
mysql> show databases;  
+-----+  
| Database |  
+-----+  
| information_schema |  
| hivemetastore |  
| movielens |  
| mysql |  
| policedb |  
| training |  
+-----+  
6 rows in set (0.01 sec)
```

Step ii. Create a new database with name *rapedb*

Step iii. Create an empty table with name *rapetb*

- ~ Use the previous created database i.e. *rapetb*

```
mysql> create database rapedb;  
mysql> use rapedb;  
mysql> create table rapetb(AREA_NAME  
varchar(30),YEAR int,RAPE_CASES_REPORTED
```

```
int,VICTIMS_BETWEEN_10_14_YRS  
int,VICTIMS_BETWEEN_14_18_YRS  
int,VICTIMS_BETWEEN_18_30_YRS  
int,VICTIMS_BETWEEN_30_50_YRS  
int,VICTIMS_ABOVE_50_YRS  
int,VICTIMS_OF_RAPE_TOTAL int);  
mysql> quit
```

```
mysql> create database rapedb;  
Query OK, 1 row affected (0.00 sec)
```

```
mysql> use rapedb;  
Database changed  
mysql> create table rapetb(AREA_NAME varchar(30),YEAR int,RAPE_CASES_REPORTED int,VICTIMS_BETWEEN_10_14_YRS int,VICTIMS_BETWEEN_14_18_YRS int,VICTIMS_BETWEEN_18_30_YRS int,VICTIMS_BETWEEN_30_50_YRS int,VICTIMS_ABOVE_50_YRS int,VICTIMS_OF_RAPE_TOTAL int);  
Query OK, 0 rows affected (0.06 sec)
```

```
mysql> quit  
Bye
```

Step iv. Now exit mysql. Export the data present on hadoop(dataset) using SQOOP into MYSQL in table *rapetb*

```
[training@localhost ~]$ hadoop fs -ls  
/user/hive/warehouse/crimeData.db/rapetb;  
[training@localhost ~]$ sqoop export \  
> --connect jdbc:mysql://localhost/rapedb \  
> --username root \  
> --table rapetb \  
> --export-dir /user/hive/warehouse/crimeData.db/rapetb \  
> --m 1;
```

```
training@localhost:~
```

```
File Edit View Terminal Tabs Help
```

```
[training@localhost ~]$ hadoop fs -ls /user/hive/warehouse/crimeData.db/rapetb;
Found 1 items
-rw-r--r-- 1 training supergroup      53231 2020-07-17 03:20 /user/hive/warehouse/crimeData.db/rapetb/RapeVictims.csv
[training@localhost ~]$ sqoop export \
> --connect jdbc:mysql://localhost/rapedb \
> --username root \
> --table rapetb \
> --export-dir /user/hive/warehouse/crimeData.db/rapetb \
> --m 1;
20/07/17 03:33:00 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset.
20/07/17 03:33:00 INFO tool.CodeGenTool: Beginning code generation
20/07/17 03:33:00 INFO manager.SqlManager: Executing SQL statement: SELECT t.* F
ROM `rapetb` AS t LIMIT 1

20/07/17 03:33:15 INFO mapred.JobClient:      Spilled Records=0
20/07/17 03:33:15 INFO mapred.JobClient:      Map output records=1470
20/07/17 03:33:15 INFO mapred.JobClient:      SPLIT_RAW_BYTES=148
20/07/17 03:33:15 INFO mapreduce.ExportJobBase: Transferred 52.1338 KB in 13.001
8 seconds (4.0097 KB/sec)
20/07/17 03:33:15 INFO mapreduce.ExportJobBase: Exported 1470 records.
```

Step v. Open mysql and use same database i.e. *rapedb*

Step vi. Perform SUM function over the data present in *rapetb*

```
[training@localhost ~]$ mysql -u root
mysql> use rapedb;
mysql> select * from rapetb limit 10;
mysql> select sum(VICTIMS_BETWEEN_10_14_YRS),
sum(VICTIMS_BETWEEN_14_18_YRS),
sum(VICTIMS_BETWEEN_18_30_YRS),
sum(VICTIMS_BETWEEN_30_50_YRS),
sum(VICTIMS_ABOVE_50_YRS) from rapetb;
```

```

[training@localhost ~]$ mysql -u root
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 30
Server version: 5.0.77 Source distribution

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.

mysql> use rapedb;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

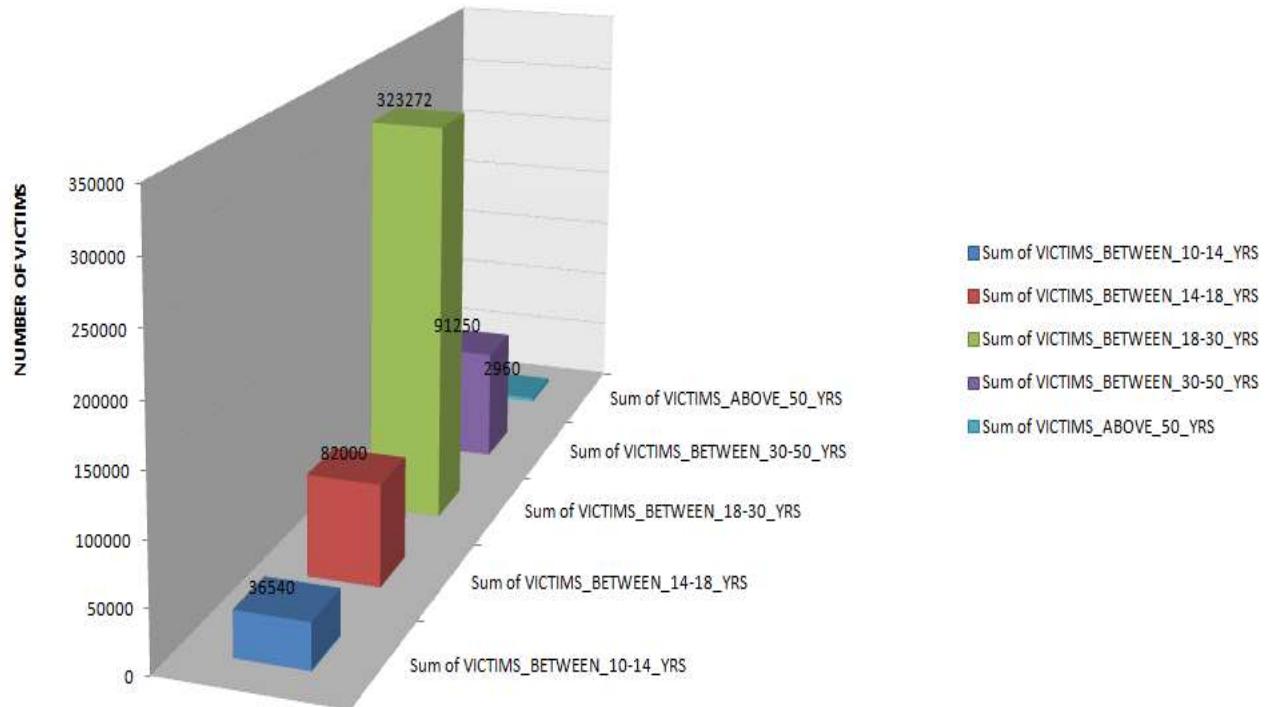
Database changed
mysql> select * from rapetb limit 10;
+-----+-----+-----+
| AREA_NAME | YEAR | RAPE_CASES_REPORTED | VICTIMS_BETWEEN_10_14_YRS | VICTIMS_BETWEEN_14_18_YRS | VICTIMS_BETWEEN_18_30_YRS | VICTIMS_BETWEEN_30_50_YRS | VICTIMS_ABOVE_50_YRS | VICTIMS_OF_RAPE_TOTAL |
+-----+-----+-----+
| sum(VICTIMS_BETWEEN_10_14_YRS) | sum(VICTIMS_BETWEEN_14_18_YRS) | sum(VICTIMS_BETWEEN_18_30_YRS) | sum(VICTIMS_BETWEEN_30_50_YRS) | sum(VICTIMS_ABOVE_50_YRS) |
|-----+-----+-----+
| 323272 | 36540 | 91250 | 82000 | 2960 |
+-----+-----+-----+
1 row in set (0.01 sec)

```

VISUALIZATION

This column graph shows the graph plotted between VICTIMS OF VARIOUS AGE GROUPS to NUMBER OF VICTIMS.

AGE WISE VICTIMS CHART GRAPH



Query 3:- Top 3 categories of Policemen killed or injured.

CODING

Step i. Open Pig in MapReduce mode

Step ii. Create a bag with name ‘a’ and load the data present on Hadoop (dataset) into bag ‘a’

```
[training@localhost ~]$ pig
grunt> a = load
'/user/hive/warehouse/crimeData.db/policekilled' using
PigStorage(',') as
(AREA_NAME:chararray, YEAR:int, CATEGORY:chararray, P
OLICE_INJURED_BY_CRIMINALS:int, POLICE_INJURED_
BY RIOTOUS_MOBS:int, POLICE_INJURED_IN_ACCIDEN
TS:int, POLICE_INJURED_IN_DACOITY_OPERATIONS:O
THER_RAIDS:int, POLICE_INJURED_IN_TERRORISTSEXT
REMISTS_OPERATIONS:int, POLICE_INJURED_ON_BOR
DER_DUTIES:int, TOTAL_POLICEMEN_INJURED:int, PO
LICE_KILLED_BY_CRIMINALS:int, POLICE_KILLED_BY_RIO
TOUS_MOBS:int, POLICE_KILLED_IN_ACCIDENTS:int, PO
LICE_KILLED_IN_DACOITY_OPERATIONS:OTHER_RAID
S:int, POLICE_KILLED_IN_TERRORISTSEXTREMISTS_OP
ERATIONS:int, POLICE_KILLED_ON_BORDER_DUTIES:int
, TOTAL_POLICEMEN_KILLED:int);
```

```
training@localhost:~$ pig
2020-07-18 09:01:57,318 [main] INFO org.apache.pig.Main - Logging error messages to: /home/training/pig_1595088117314.log
2020-07-18 09:01:57,746 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to hadoop file system at: hdfs://localhost:8020
2020-07-18 09:01:58,126 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to map-reduce job tracker at: localhost:8021
grunt> a = load '/user/hive/warehouse/crimeData.db/policekilled' using PigStorage(',') as (AREA_NAME:chararray, YEAR:int, CATEGORY:chararray, POLICE_INJURED_BY_CRIMINALS:int, POLICE_INJURED_BY RIOTOUS_MOBS:int, POLICE_INJURED_IN_ACCIDENTS:int, POLICE_INJURED_IN_DAOITY_OPERATIONSOTHER_RAIDS:int, POLICE_INJURED_IN_TERRORISTSEX_TREMISTS_OPERATIONS:int, POLICE_INJURED_ON_BORDER_DUTIES:int, TOTAL_POLICEMEN_INJURED:int, POLICE_KILLED_BY_CRIMINALS:int, POLICE_KILLED_BY RIOTOUS_MOBS:int, POLICE_KILLED_IN_ACCIDENTS:int, POLICE_KILLED_IN_DAOITY_OPERATIONSOTHER_RAIDS:int, POLICE_KILLED_IN_TERRORISTSEX_TREMISTS_OPERATIONS:int, POLICE_KILLED_ON_BORDER_DUTIES:int, TOTAL_POLICEMEN_KILLED:int);
```

Step iii. Now perform GROUP operation by referencing with a column CATEGORY in a bag '**b**'

Step iv. Perform COUNT operation on '**a**' in bag '**c**'

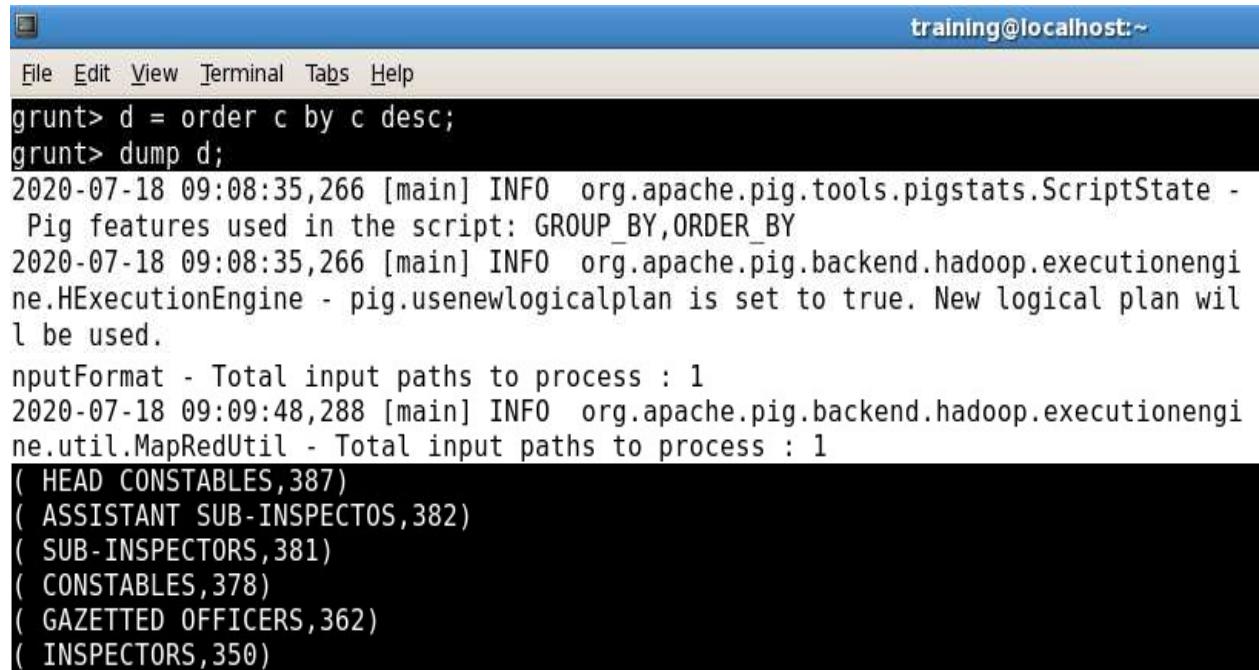
```
grunt> b = group a by CATEGORY;
grunt> c = foreach b generate group,COUNT(a) as c;
grunt> dump c;
```

```
training@localhost:~$ pig
File Edit View Terminal Tabs Help
grunt> b = group a by CATEGORY;
grunt> c = foreach b generate group,COUNT(a) as c;
grunt> dump c;
2020-07-18 09:06:16,416 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: GROUP_BY
2020-07-18 09:06:16,416 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - pig.usenewlogicalplan is set to true. New logical plan will be used.
```

```
2020-07-18 09:06:41,665 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-07-18 09:06:41,665 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
( CONSTABLES,378)
( INSPECTORS,350)
( SUB-INSPECTORS,381)
( HEAD CONSTABLES,387)
( GAZETTED OFFICERS,362)
( ASSISTANT SUB-INSPECTOS,382)
```

Step v. Order bag ‘c’ as Descending Order

```
grunt> d = order c by c desc;
grunt> dump d;
```



A screenshot of a terminal window titled "training@localhost:~". The window shows the following command history and output:

```
File Edit View Terminal Tabs Help
grunt> d = order c by c desc;
grunt> dump d;
2020-07-18 09:08:35,266 [main] INFO org.apache.pig.tools.pigstats.ScriptState -
  Pig features used in the script: GROUP_BY,ORDER_BY
2020-07-18 09:08:35,266 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine -
  pig.usenewlogicalplan is set to true. New logical plan will be used.
inputFormat - Total input paths to process : 1
2020-07-18 09:09:48,288 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
( HEAD CONSTABLES,387)
( ASSISTANT SUB-INSPECTOS,382)
( SUB-INSPECTORS,381)
( CONSTABLES,378)
( GAZETTED OFFICERS,362)
( INSPECTORS,350)
```

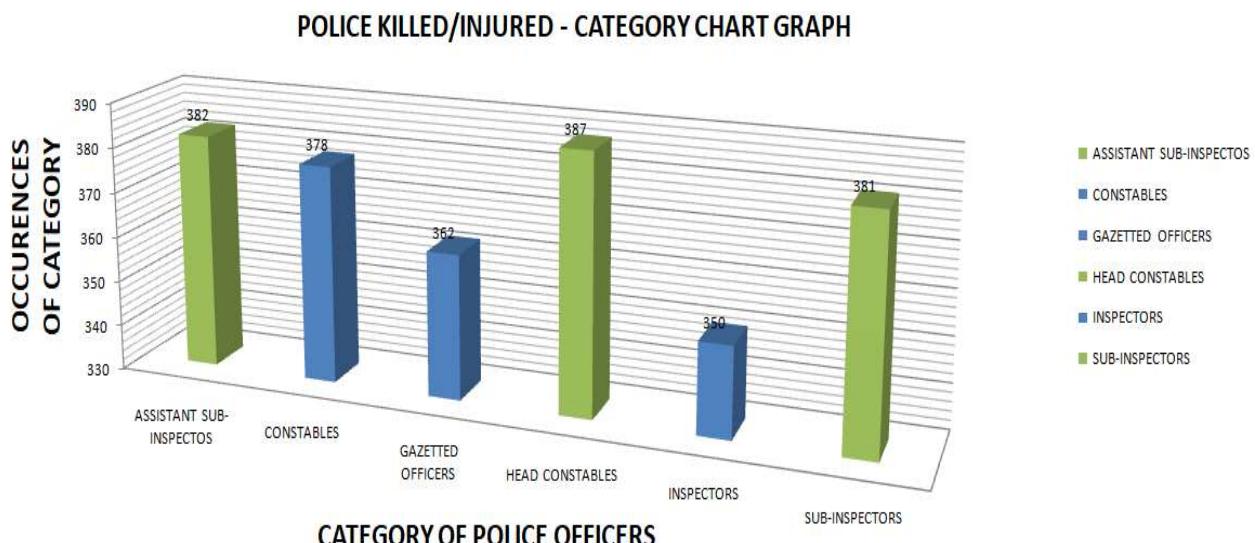
Step vi. Apply LIMIT operation as limit 3.

```
grunt> e = limit d 3;  
grunt> dump e;
```

```
File Edit View Terminal Tabs Help  
training@localhost:~  
grunt> e = limit d 3;  
grunt> dump e;  
2020-07-18 09:11:34,367 [main] INFO org.apache.pig.tools.pigstats.ScriptState -  
Pig features used in the script: GROUP_BY,ORDER_BY,LIMIT  
2020-07-18 09:11:34,367 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - pig.usenewlogicalplan is set to true. New logical plan will use.mapReduceLayer.MapReduceLauncher - Success!  
2020-07-18 09:12:42,707 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1  
2020-07-18 09:12:42,707 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1  
( HEAD CONSTABLES,387)  
( ASSISTANT SUB-INSPECTOS,382)  
( SUB-INSPECTORS,381)
```

VISUALIZATION

This column graph shows the graph plotted between CATEGORY OF POLICE OFFICERS and their OCCURENCES. The **THREE BAR** with **GREEN COLOR** shows the top 3 category in the graph.



Query 4:- Categorization of Murder victims on the basis of Gender.

CODING

Step i. Open Hive and use previously created database i.e.

crimedata

Step ii. Create a new table with name *murdervictim*

```
[training@localhost ~]$ hive
hive> use crimedata;
hive> create table murdervictimtb(AREA_NAME string,YEAR
int,GENDER string,VICTIMS_UPTO_10_YRS
int,VICTIMS_UPTO_10_15_YRS
int,VICTIMS_UPTO_15_18_YRS
int,VICTIMS_UPTO_18_30_YRS
int,VICTIMS_UPTO_30_50_YRS
int,VICTIMS_ABOVE_50_YRS int,VICTIMS_TOTAL int)
> row format delimited
> fields terminated by ','
> lines terminated by '\n'
> stored as textfile;
```

```
training@localhost:~
```

```
File Edit View Terminal Tabs Help
[training@localhost ~]$ hive
Hive history file=/tmp/training/hive_job_log_training_202007172232_1758519624.tx
t
hive> show databases;
OK
crimedata
default
Time taken: 2.982 seconds
hive> use crimedata;
OK
Time taken: 0.025 seconds
hive> create table murdervictimtb(AREA_NAME string, YEAR int, GENDER string, VICTIM
S upto_10_YRS int, VICTIMS upto_10_15_YRS int, VICTIMS upto_15_18_YRS int, VICTIMS
upto_18_30_YRS int, VICTIMS upto_30_50_YRS int, VICTIMS above_50_YRS int, VICTIMS_T
OTAL int)
    > row format delimited
    > fields terminated by ','
    > lines terminated by '\n'
    > stored as textfile;
OK
Time taken: 0.747 seconds
```

Step iii. Load data into the table from local file system (location of dataset)

```
hive> load data local inpath
'/home/training/IBMPProject/MurderVictim.csv' into table
murdervictimtb;
```

```
hive> load data local inpath '/home/training/IBMPProject/MurderVictim.csv' into t
able murdervictimtb;
Copying data from file:/home/training/IBMPProject/MurderVictim.csv
Copying file: file:/home/training/IBMPProject/MurderVictim.csv
Loading data to table crimedata.murdervictimtb
OK
Time taken: 0.541 seconds
```

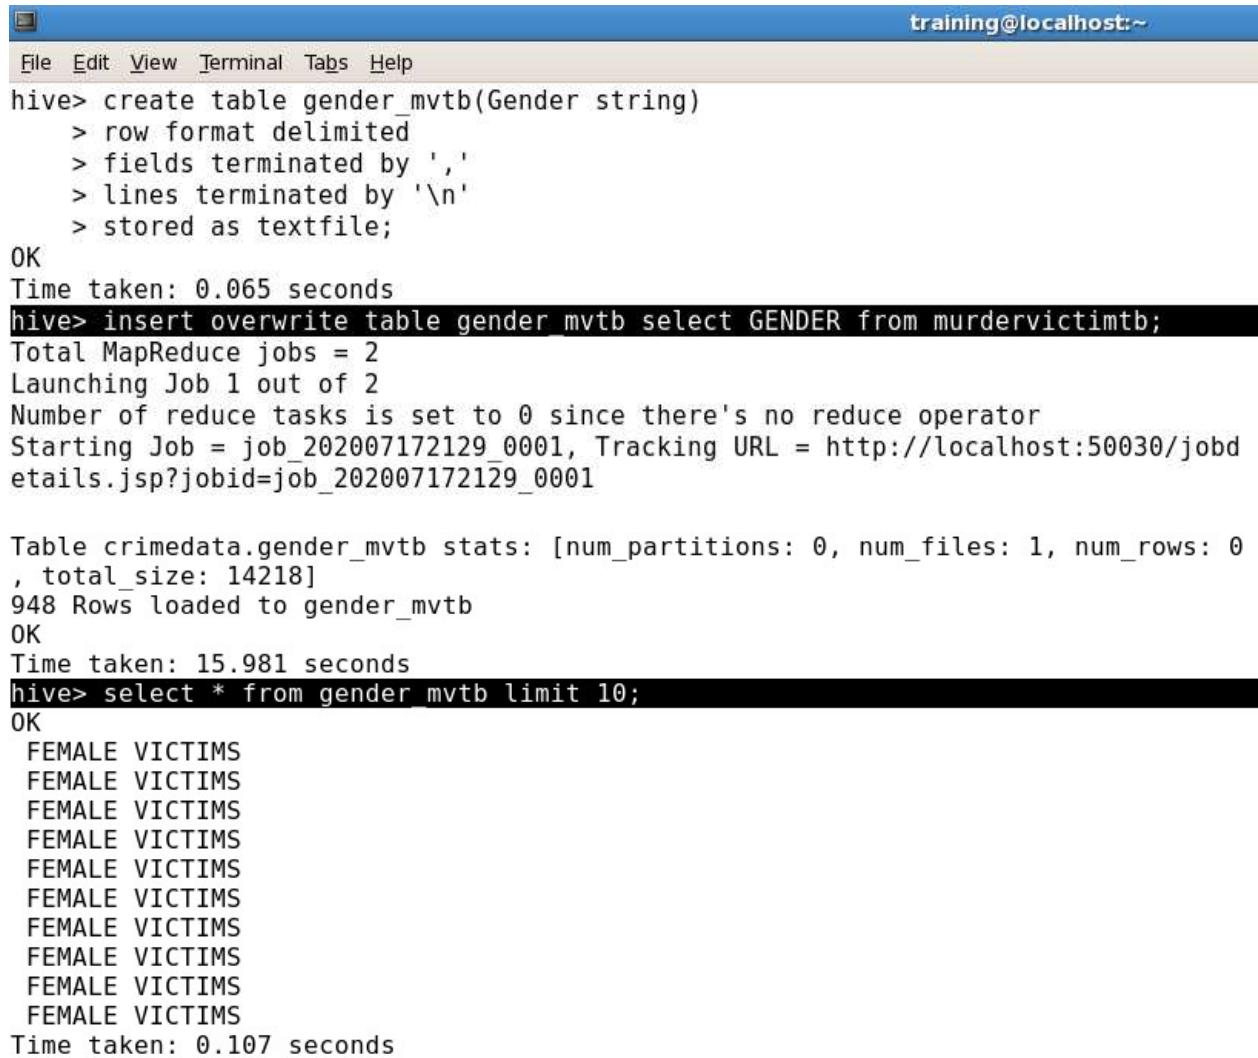
Step iv. Create a new table with name *gender_mvtb*

Step v. Load data of only GENDER column from table
murdervictim

```

hive> create table gender_mvtb(Gender string)
      > row format delimited
      > fields terminated by ','
      > lines terminated by '\n'
      > stored as textfile;
hive> insert overwrite table gender_mvtb select GENDER
      from murdervictimtb;

```



The screenshot shows a terminal window with a blue header bar containing the text "training@localhost:~". Below the header is a menu bar with "File", "Edit", "View", "Terminal", "Tabs", and "Help". The main area of the terminal displays the following Hive session:

```

hive> create table gender_mvtb(Gender string)
      > row format delimited
      > fields terminated by ','
      > lines terminated by '\n'
      > stored as textfile;
OK
Time taken: 0.065 seconds
hive> insert overwrite table gender_mvtb select GENDER from murdervictimtb;
Total MapReduce jobs = 2
Launching Job 1 out of 2
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_202007172129_0001, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007172129_0001

Table crimedata.gender_mvtb stats: [num_partitions: 0, num_files: 1, num_rows: 0
, total_size: 14218]
948 Rows loaded to gender_mvtb
OK
Time taken: 15.981 seconds
hive> select * from gender_mvtb limit 10;
OK
FEMALE VICTIMS
Time taken: 0.107 seconds

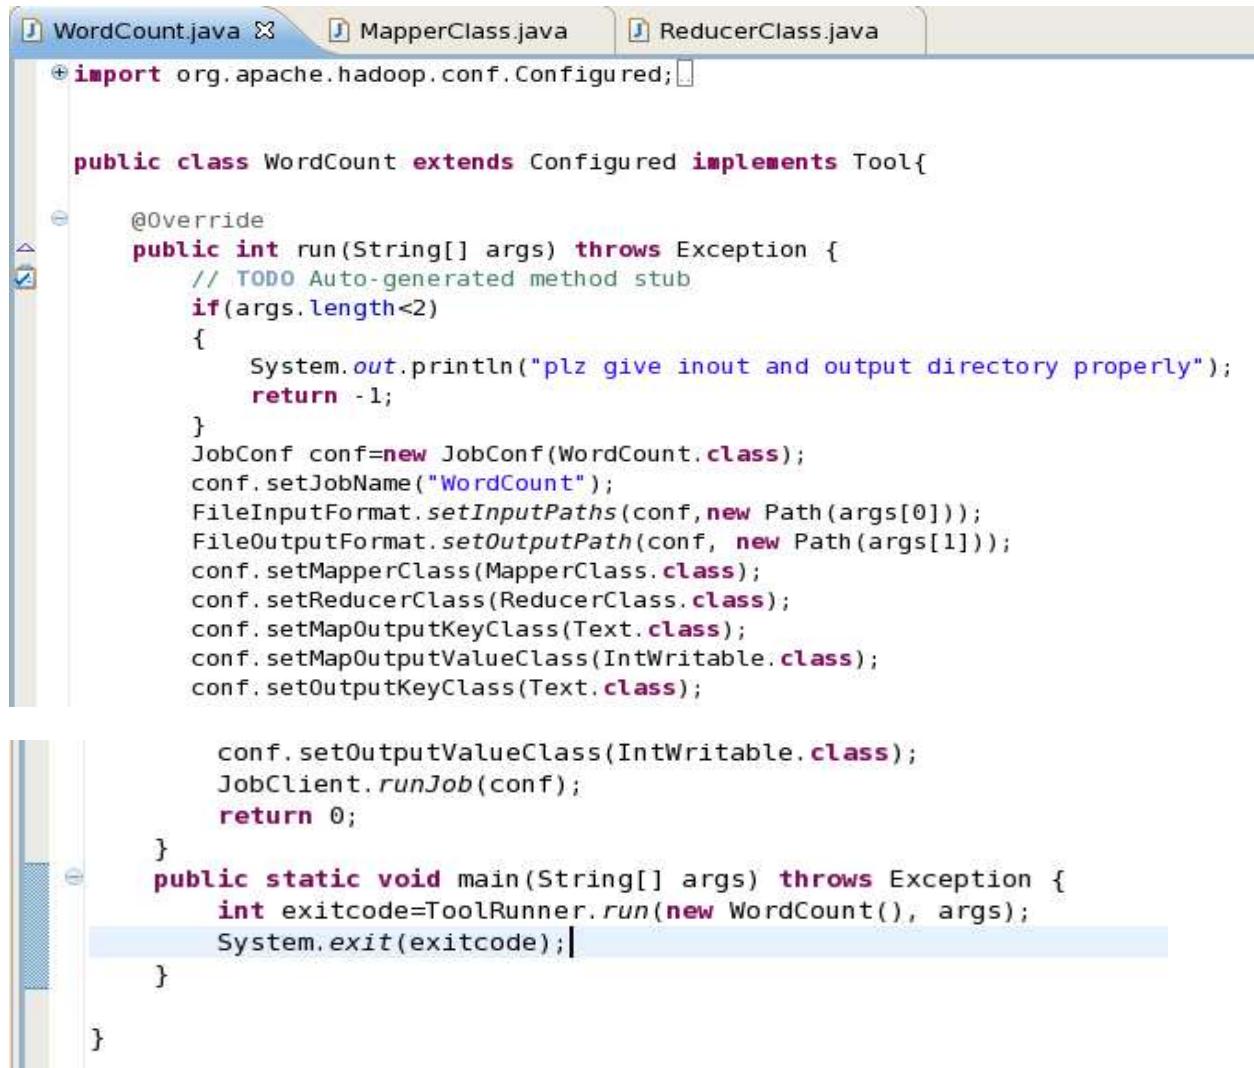
```

Step vi. Now minimize the terminal and open ECLIPSE

Step vii. Create a new java project with name *WordCount*.

- ~ The project contains three classes i.e. WordCount, Mapper, Reducer Class along with codes for MapReduce.

Screenshot: WordCount Class



The screenshot shows a Java code editor with three tabs at the top: WordCount.java, MapperClass.java, and ReducerClass.java. The WordCount.java tab is active, displaying the following code:

```
import org.apache.hadoop.conf.Configuration;

public class WordCount extends Configuration implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        // TODO Auto-generated method stub
        if(args.length<2)
        {
            System.out.println("plz give inout and output directory properly");
            return -1;
        }
        JobConf conf=new JobConf(WordCount.class);
        conf.setJobName("WordCount");
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
        conf.setMapperClass(MapperClass.class);
        conf.setReducerClass(ReducerClass.class);
        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);

        conf.setOutputValueClass(IntWritable.class);
        JobClient.runJob(conf);
        return 0;
    }
    public static void main(String[] args) throws Exception {
        int exitcode=ToolRunner.run(new WordCount(), args);
        System.exit(exitcode);
    }
}
```

Screenshot: Mapper Class

The screenshot shows a Java code editor with three tabs at the top: WordCount.java, MapperClass.java (which is currently selected), and ReducerClass.java. The MapperClass.java tab has a blue checkmark icon. The code in the editor is:

```
public class MapperClass extends MapReduceBase implements Mapper<LongWritable,Text,Text,IntWritable>{

    @Override
    public void map(LongWritable key, Text value,
                    OutputCollector<Text, IntWritable> output, Reporter arg3)
        throws IOException {
        // TODO Auto-generated method stub
        String s=value.toString();
        for(String word:s.split(" "))
        {
            if(word.length()>0)
            {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }

}
```

Screenshot: Reducer Class

The screenshot shows a Java code editor with three tabs at the top: WordCount.java, MapperClass.java, and ReducerClass.java (which is currently selected). The ReducerClass.java tab has a blue checkmark icon. The code in the editor is:

```
import java.io.IOException;
import java.util.Iterator;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class ReducerClass extends MapReduceBase implements Reducer<Text,IntWritable,Text,IntWritable>{

    @Override
    public void reduce(Text key, Iterator<IntWritable> value,
                      OutputCollector<Text, IntWritable> output, Reporter arg3)
        throws IOException {
        // TODO Auto-generated method stub
        int count=0;
        while(value.hasNext())
        {
            IntWritable i=value.next();

            count+=i.get();
        }
        output.collect(key,new IntWritable(count));
    }

}
```

Step viii. Export that project and create a jar file with name *wordcount.jar*

Step ix. Use MapReduce commands in terminal to perform wordcount over the data present in gender_mvtb

```
[training@localhost ~]$ hadoop jar  
/home/training/Desktop/wordcount.jar WordCount  
/user/hive/warehouse/crimeData.db/gender_mvtb  
wordcountoutput9  
[training@localhost ~]$ hadoop fs -cat  
/user/training/wordcountoutput9/part-00000
```

The screenshot shows a terminal window with a blue header bar containing the text "training@localhost:~". Below the header is a menu bar with "File", "Edit", "View", "Terminal", "Tabs", and "Help". The main area of the terminal shows the command-line interaction:

```
File Edit View Terminal Tabs Help  
[training@localhost ~]$ hadoop jar /home/training/Desktop/wordcount.jar WordCount  
/user/hive/warehouse/crimeData.db/gender_mvtb wordcountoutput9  
20/07/18 04:34:36 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should implement Tool for the same.  
20/07/18 04:34:36 WARN snappy.LoadSnappy: Snappy native library is available  
20/07/18 04:34:36 INFO util.NativeCodeLoader: Loaded the native-hadoop library  
20/07/18 04:34:36 INFO snappy.LoadSnappy: Snappy native library loaded  
[training@localhost ~]$ hadoop fs -cat /user/training/wordcountoutput9/part-00000  
FEMALE 473  
MALE 475  
VICTIMS 948
```

VISUALIZATION

This pie-chart graph shows the graph plotted between numbers of MALE & FEMALE VICTIMS in murder crime.

MALE - FEMALE CHART GRAPH OF MURDER VICTIMS

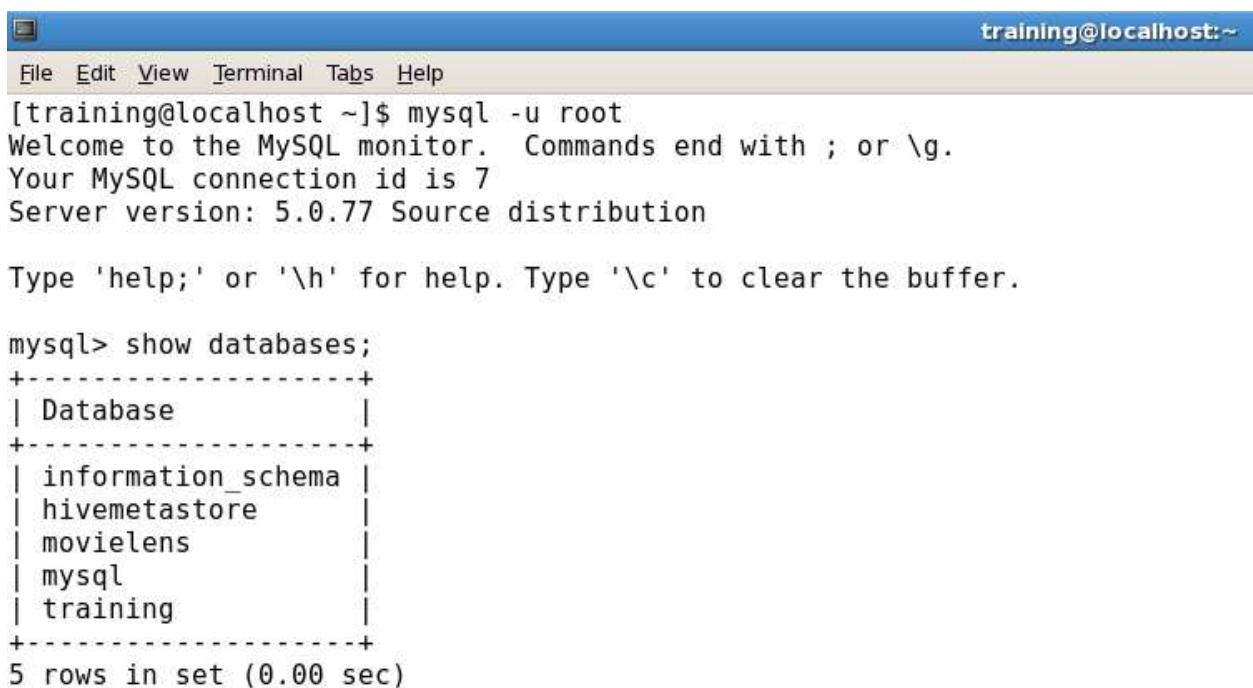


Query 5:- Year having maximum number of casualties (Policemen killed).

CODING

Step i. Open Mysql

```
[training@localhost ~]$ mysql -u root  
mysql> show databases;
```



A screenshot of a terminal window titled "training@localhost:~". The window shows the MySQL monitor starting up, displaying connection details and the server version. It then executes the "show databases;" command, listing five databases: information_schema, hivemetastore, movielens, mysql, and training. The output concludes with "5 rows in set (0.00 sec)".

```
File Edit View Terminal Tabs Help  
[training@localhost ~]$ mysql -u root  
Welcome to the MySQL monitor. Commands end with ; or \g.  
Your MySQL connection id is 7  
Server version: 5.0.77 Source distribution  
  
Type 'help;' or '\h' for help. Type '\c' to clear the buffer.  
  
mysql> show databases;  
+-----+  
| Database      |  
+-----+  
| information_schema |  
| hivemetastore   |  
| movielens       |  
| mysql           |  
| training         |  
+-----+  
5 rows in set (0.00 sec)
```

Step ii. Create a new database with name *policedb*

Step iii. Create an empty new table with name *policetb* by using previously created database i.e. *policedb*

```
mysql> create database policedb;  
mysql> use policedb;  
mysql> create table policetb(AREA_NAME  
varchar(30),YEAR int,CATEGORY  
varchar(30),POLICE_INJURED_BY_CRIMINALS
```

```
int,POLICE_INJURED_BY_RIOTOUS_MOBS  
int,POLICE_INJURED_IN_ACCIDENTS  
int,POLICE_INJURED_IN_DACOITY_OPERATIONSOTHER_  
R_RAIDS  
int,POLICE_INJURED_IN_TERRORISTSEXTEMMISTS_OP  
ERATIONS int,POLICE_INJURED_ON_BORDER_DUTIES  
int,TOTAL_POLICEMEN_INJURED  
int,POLICE_KILLED_BY_CRIMINALS  
int,POLICE_KILLED_BY_RIOTOUS_MOBS  
int,POLICE_KILLED_IN_ACCIDENTS  
int,POLICE_KILLED_IN_DACOITY_OPERATIONSOTHER_  
RAIDS  
int,POLICE_KILLED_IN_TERRORISTSEXTEMMISTS_OPE  
RATIONS int,POLICE_KILLED_ON_BORDER_DUTIES  
int,TOTAL_POLICEMEN_KILLED int);  
mysql> quit
```

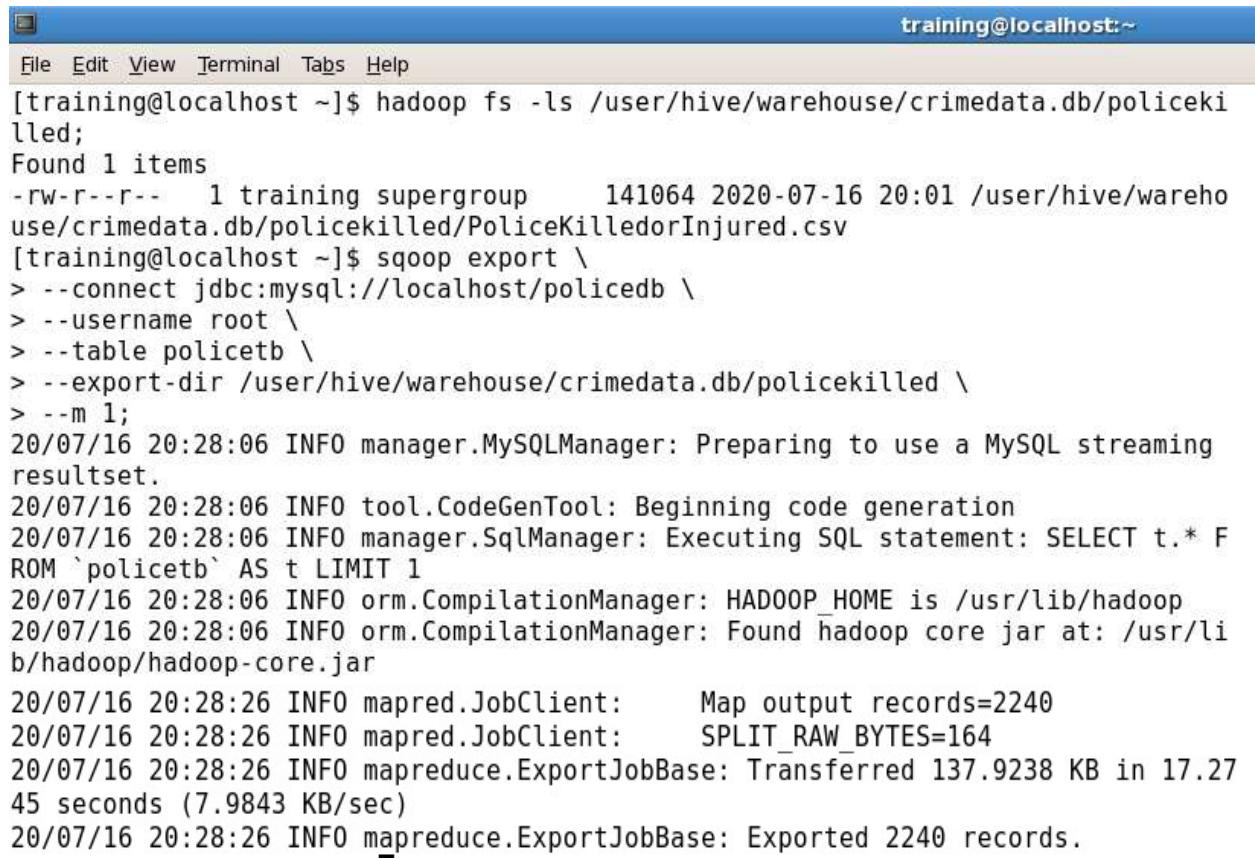
```
mysql> create database policedb;  
Query OK, 1 row affected (0.00 sec)
```

```
mysql> use policedb;  
Database changed  
mysql> create table policetb(AREA_NAME varchar(30),YEAR int,CATEGORY varchar(30)  
,POLICE_INJURED_BY_CRIMINALS int,POLICE_INJURED_BY_RIOTOUS_MOBS int,POLICE_INJUR  
ED_IN_ACCIDENTS int,POLICE_INJURED_IN_DACOITY_OPERATIONSOTHER_RAIDS int,POLICE_I  
NJURED_IN_TERRORISTSEXTEMMISTS_OPERATIONS int,POLICE_INJURED_ON_BORDER_DUTIES in  
t,TOTAL_POLICEMEN_INJURED int,POLICE_KILLED_BY_CRIMINALS int,POLICE_KILLED_BY_RI  
OTOUS_MOBS int,POLICE_KILLED_IN_ACCIDENTS int,POLICE_KILLED_IN_DACOITY_OPERATION  
SOTHER_RAIDS int,POLICE_KILLED_IN_TERRORISTSEXTEMMISTS_OPERATIONS int,POLICE_KIL  
LED_ON_BORDER_DUTIES int,TOTAL_POLICEMEN_KILLED int);  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> quit  
Bye
```

Step iv. Now exit mysql. Export the data present on hadoop(dataset) using SQOOP into MYSQL in table *policetb*

```
[training@localhost ~]$ hadoop fs -ls
/usr/hive/warehouse/crимedata.db/policekilled;
[training@localhost ~]$ sqoop export \
> --connect jdbc:mysql://localhost/policedb \
> --username root \
> --table policetb \
> --export-dir /user/hive/warehouse/crимedata.db/policekilled
\
> --m 1;
```



The screenshot shows a terminal window titled "training@localhost:~". The window contains the following command history:

```
File Edit View Terminal Tabs Help
[training@localhost ~]$ hadoop fs -ls /user/hive/warehouse/crимedata.db/policekilled;
Found 1 items
-rw-r--r-- 1 training supergroup 141064 2020-07-16 20:01 /user/hive/warehouse/crимedata.db/policekilled/PoliceKilledorInjured.csv
[training@localhost ~]$ sqoop export \
> --connect jdbc:mysql://localhost/policedb \
> --username root \
> --table policetb \
> --export-dir /user/hive/warehouse/crимedata.db/policekilled \
> --m 1;
20/07/16 20:28:06 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
20/07/16 20:28:06 INFO tool.CodeGenTool: Beginning code generation
20/07/16 20:28:06 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `policetb` AS t LIMIT 1
20/07/16 20:28:06 INFO orm.CompilationManager: HADOOP_HOME is /usr/lib/hadoop
20/07/16 20:28:06 INFO orm.CompilationManager: Found hadoop core jar at: /usr/lib/hadoop/hadoop-core.jar
20/07/16 20:28:26 INFO mapred.JobClient: Map output records=2240
20/07/16 20:28:26 INFO mapred.JobClient: SPLIT_RAW_BYTES=164
20/07/16 20:28:26 INFO mapreduce.ExportJobBase: Transferred 137.9238 KB in 17.2745 seconds (7.9843 KB/sec)
20/07/16 20:28:26 INFO mapreduce.ExportJobBase: Exported 2240 records.
```

Step v. Open mysql and use same database i.e. *policedb*

Step vi. Perform MAX function over the data present in *policetb*

```
[training@localhost ~]$ mysql -u root
mysql> use policedb;
mysql> select * from policetb limit 10;
mysql> select YEAR,MAX(TOTAL_POLICEMEN_KILLED)
from policetb;
```

```
File Edit View Terminal Tabs Help
[training@localhost ~]$ mysql -u root
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 11
Server version: 5.0.77 Source distribution

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.
```

```
mysql> use policedb;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

```
Database changed
mysql> select * from policetb limit 10;
+-----+-----+-----+
|-----+-----+-----+
|-----+-----+-----+
|-----+-----+-----+
```

```
File Edit View Terminal Tabs Help
[training@localhost ~]$ mysql -u root
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 20
Server version: 5.0.77 Source distribution

Type 'help;' or '\h' for help. Type '\c' to clear the buffer.
```

```
mysql> use policedb;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

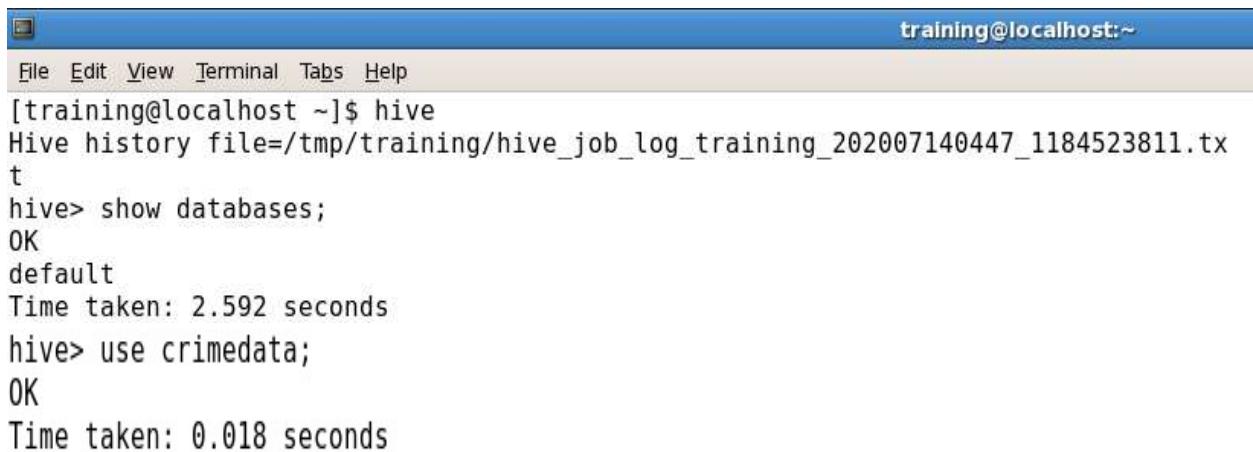
```
Database changed
mysql> select YEAR,max(TOTAL_POLICEMEN_KILLED) from policetb;
+-----+-----+
| YEAR | max(TOTAL_POLICEMEN_KILLED) |
+-----+-----+
| 2001 | 116 |
+-----+-----+
1 row in set (0.01 sec)
```

Query 6:- Total number of Dacoity happened in each year.

CODING

Step i. Open hive and use previously created database i.e. crimedata

```
[training@localhost ~]$ hive  
hive> show databases;  
hive> use crimedata;
```



A screenshot of a terminal window titled "training@localhost:~". The window shows the following command history:

```
File Edit View Terminal Tabs Help  
[training@localhost ~]$ hive  
Hive history file=/tmp/training/hive_job_log_training_202007140447_1184523811.tx  
t  
hive> show databases;  
OK  
default  
Time taken: 2.592 seconds  
hive> use crimedata;  
OK  
Time taken: 0.018 seconds
```

Step ii. Perform SUM function over the column DACOITY by grouping YEAR

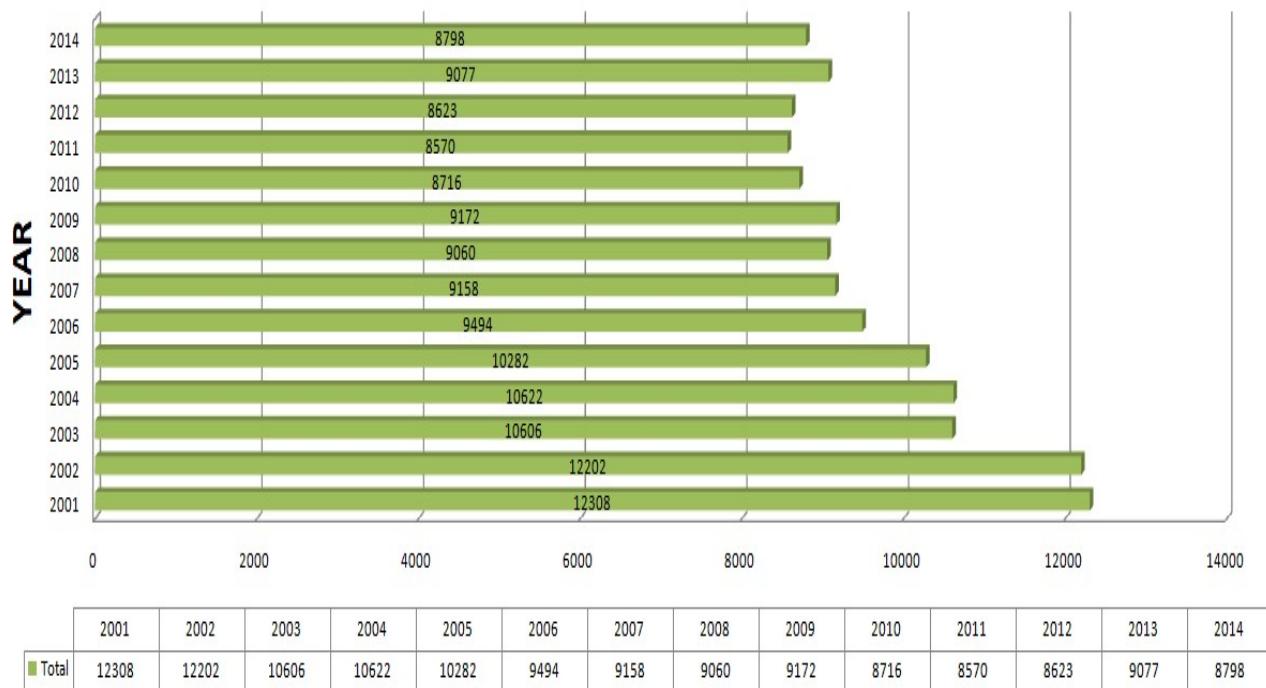
```
hive> select YEAR,sum(DACOITY) from crimetable group by  
YEAR;
```

```
File Edit View Terminal Tabs Help
hive> select YEAR,sum(DACOITY) from crimetble group by YEAR;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202007140415_0019, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007140415_0019
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007140415_0019
2020-07-15 19:35:02,656 Stage-1 map = 0%,  reduce = 0%
2020-07-15 19:35:04,687 Stage-1 map = 100%,  reduce = 0%
2020-07-15 19:35:11,974 Stage-1 map = 100%,  reduce = 33%
2020-07-15 19:35:12,984 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_202007140415_0019
OK
NULL      NULL
2001      12308
2002      12202
2003      10606
2004      10622
2005      10282
2006      9494
2007      9158
2008      9060
2009      9172
2010      8716
2011      8570
2012      8623
2013      9077
2014      8798
Time taken: 14.538 seconds
```

VISUALIZATION

This bar graph shows the graph plotted between DACOITY happened in each YEAR.

DACITY CHART GRAPH

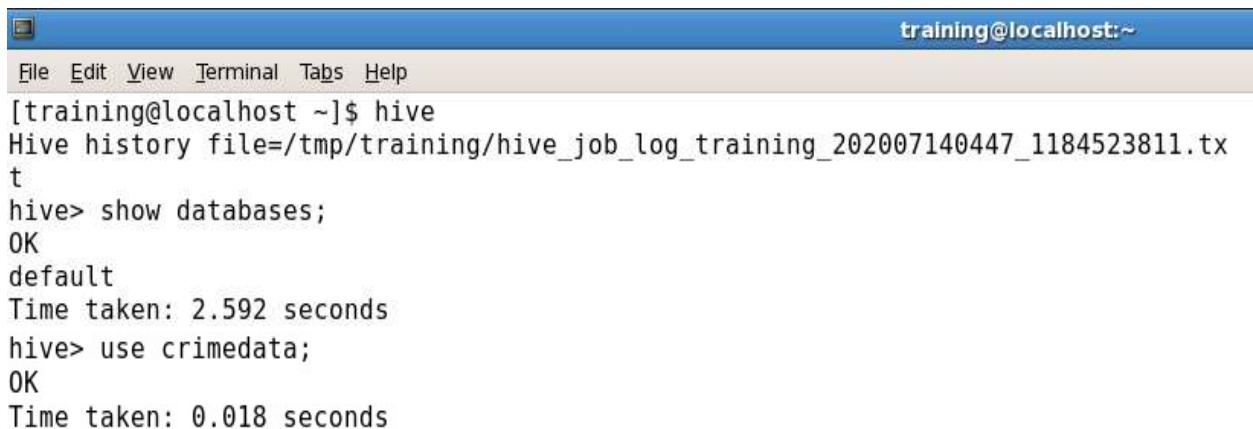


Query 7:- Top 5 states of Dowry Deaths during the year 2014.

CODING

Step i. Open hive and use previously created database i.e. crimedata

```
[training@localhost ~]$ hive  
hive> show databases;  
hive> use crimedata;
```



```
File Edit View Terminal Tabs Help  
[training@localhost ~]$ hive  
Hive history file=/tmp/training/hive_job_log_training_202007140447_1184523811.tx  
t  
hive> show databases;  
OK  
default  
Time taken: 2.592 seconds  
hive> use crimedata;  
OK  
Time taken: 0.018 seconds
```

Step ii. Create a new table with name *dowrydeaths*

Step iii. Load the data from previously created table i.e. *crimetable*

```
hive> create table dowrydeaths(STATE_UT  
string,DOWRY_DEATHS int) row format delimited fields  
terminated by ',' lines terminated by '\n' stored as textfile;  
hive> insert overwrite table dowrydeaths select  
STATE_UT,sum(DOWRY_DEATHS) from crimetable where  
YEAR='2014' group by STATE_UT;
```

```

training@localhost:~
```

File Edit View Terminal Tabs Help

```

hive> create table dowrydeaths(STATE_UT string,DOWRY_DEATHS int) row format deli
mited fields terminated by ',' lines terminated by '\n' stored as textfile;
OK
Time taken: 0.258 seconds
```

```

training@localhost:~
```

File Edit View Terminal Tabs Help

```

hive> insert overwrite table dowrydeaths select STATE_UT,sum(DOWRY_DEATHS) from
crimetable where YEAR='2014' group by STATE_UT;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202007140415_0014, Tracking URL = http://localhost:50030/jobdet
ails.jsp?jobid=job_202007140415_0014
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:80
21 -kill job_202007140415_0014
2020-07-15 04:44:10,533 Stage-1 map = 0%,  reduce = 0%
2020-07-15 04:44:13,570 Stage-1 map = 100%,  reduce = 0%
2020-07-15 04:44:20,634 Stage-1 map = 100%,  reduce = 33%
2020-07-15 04:44:21,646 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_202007140415_0014
Loading data to table crimedata.dowrydeaths
Deleted hdfs://localhost/user/hive/warehouse/crimedata.db/dowrydeaths
Table crimedata.dowrydeaths stats: [num_partitions: 0, num_files: 1, num_rows: 0
, total_size: 491]
36 Rows loaded to dowrydeaths
OK
Time taken: 20.656 seconds
```

Step iv. Perform **desc** function with **limit 5** over the data present in dowrydeaths

```

hive> select * from dowrydeaths sort by DOWRY_DEATHS
desc limit 5;
```

```

hive> select * from dowrydeaths sort by DOWRY_DEATHS desc limit 5;
Total MapReduce jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202007140415_0016, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007140415_0016
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007140415_0016
2020-07-15 04:46:14,780 Stage-1 map = 0%,  reduce = 0%
2020-07-15 04:46:16,798 Stage-1 map = 100%,  reduce = 0%
2020-07-15 04:46:23,895 Stage-1 map = 100%,  reduce = 33%
2020-07-15 04:46:24,905 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_202007140415_0016

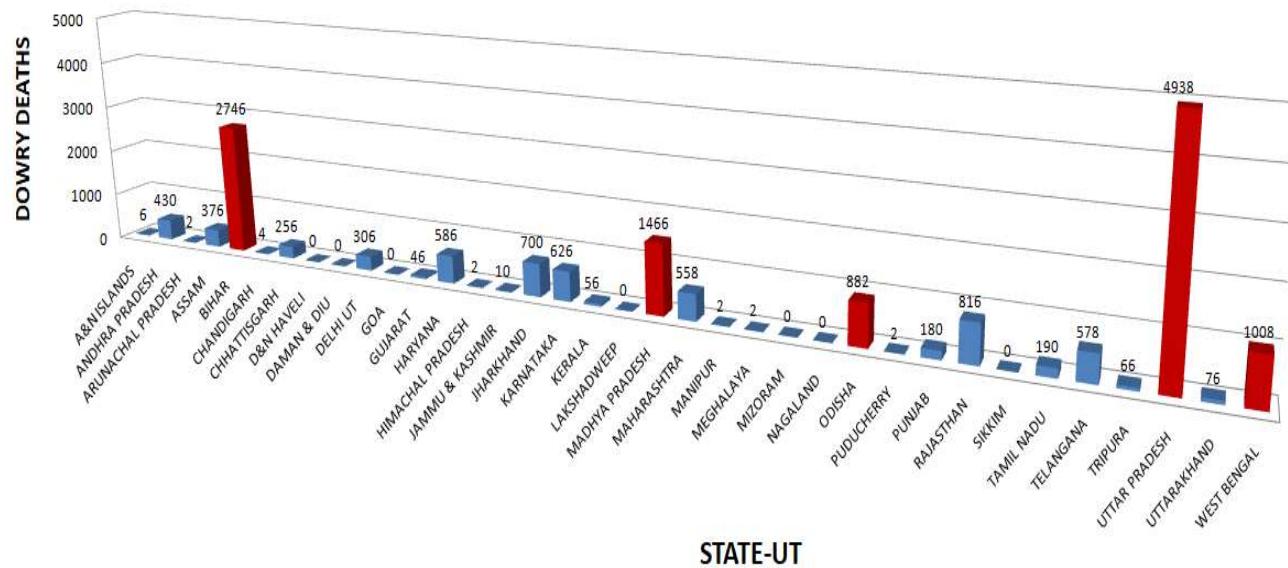
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202007140415_0017, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007140415_0017
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007140415_0017
2020-07-15 04:46:28,168 Stage-2 map = 0%,  reduce = 0%
2020-07-15 04:46:30,221 Stage-2 map = 100%,  reduce = 0%
2020-07-15 04:46:37,283 Stage-2 map = 100%,  reduce = 33%
2020-07-15 04:46:38,293 Stage-2 map = 100%,  reduce = 100%
Ended Job = job_202007140415_0017
OK
UTTAR PRADESH    4938
BIHAR      2746
MADHYA PRADESH   1466
WEST BENGAL     1008
ODISHA       882
Time taken: 26.895 seconds

```

VISUALIZATION

This column graph shows the graph plotted between STATE/UT and DOWRY DEATHS. The **FIVE BAR** with **RED COLOR** shows the top 5 category in the graph.

STATE-DOWRY DEATHS CHART GRAPH



Query 8:- Top 3 category of Auto Theft during year 2011.

CODING

- Step i. Open Pig in local mode
- Step ii. Create a bag with name ‘a’ and load the data present on Hadoop (dataset) into bag ‘a’

```
[training@localhost ~]$ pig -x local
grunt> a = load
'/home/training/IBMPProject/AutoTheft.csv' using
PigStorage(',') as
(AREA_NAME:chararray, YEAR:int, CATEGORY:chararr
ay, AUTO_THEFT_COORDINATED:int, AUTO_THEFT_
RECOVERED:int, AUTO_THEFT_STOLEN:int);
```



A screenshot of a terminal window titled "training@localhost:~". The window shows the command "pig -x local" being run, followed by the definition of a bag "a" loading the CSV file "/home/training/IBMPProject/AutoTheft.csv" using PigStorage(')'). The schema for the bag is defined with columns: AREA_NAME (chararray), YEAR (int), CATEGORY (chararray), AUTO_THEFT_COORDINATED (int), AUTO_THEFT_RECOVERED (int), and AUTO_THEFT_STOLEN (int).

```
File Edit View Terminal Tabs Help
[training@localhost ~]$ pig -x local
2020-07-18 03:04:01,544 [main] INFO org.apache.pig.Main - Logging error messages to: /home/training/pig_1595066641542.log
2020-07-18 03:04:01,685 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to hadoop file system at: file:///-
grunt> a = load '/home/training/IBMPProject/AutoTheft.csv' using PigStorage(',') as (AREA_NAME:chararray, YEAR:int, CATEGORY:chararray, AUTO_THEFT_COORDINATED:int, AUTO_THEFT_RECOVERED:int, AUTO_THEFT_STOLEN:int);
```

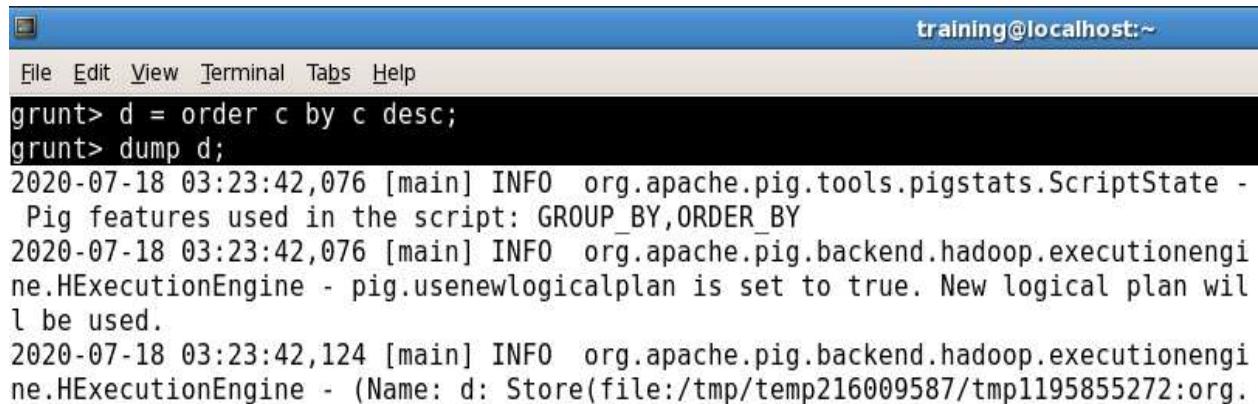
- Step iii. Now perform GROUP operation by referencing with a column CATEGORY in a bag ‘b’
- Step iv. Perform COUNT operation on ‘a’ in bag ‘c’

```
grunt> b = group a by CATEGORY;
grunt> c = foreach b generate group,COUNT(a) as c;
grunt> dump c;
```

```
grunt> b = group a by CATEGORY;
grunt> c = foreach b generate group,COUNT(a) as c;
grunt> dump c;
2020-07-18 03:15:30,940 [main] INFO org.apache.pig.tools.pigstats.ScriptState -
Pig features used in the script: GROUP_BY
2020-07-18 03:15:30,941 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - pig.usenewlogicalplan is set to true. New logical plan will be used.
2020-07-18 03:15:31,166 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - (Name: c: Store(file:/tmp/temp216009587/tmp1421318702:org.apache.pig.impl.io.InterStorage) - scope-30 Operator Key: scope-30)
nputFormat - Total input paths to process : 1
2020-07-18 03:15:40,781 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(BUSES,103)
( BUSES,254)
(TRUCKS/TEMPO ,124)
( CAR/TAXI/JEEP,465)
(OTHER VEHICLES,122)
( CYCLES/ SCOOTERS,480)
( OTHER MOTOR VEHICLES,303)
( GOODS CARRYING VEHICLES (TRUCKS/TEMPO ETC),292)
```

Step v. Order bag ‘c’ as Descending Order

```
grunt> d = order c by c desc;
grunt> dump d;
```



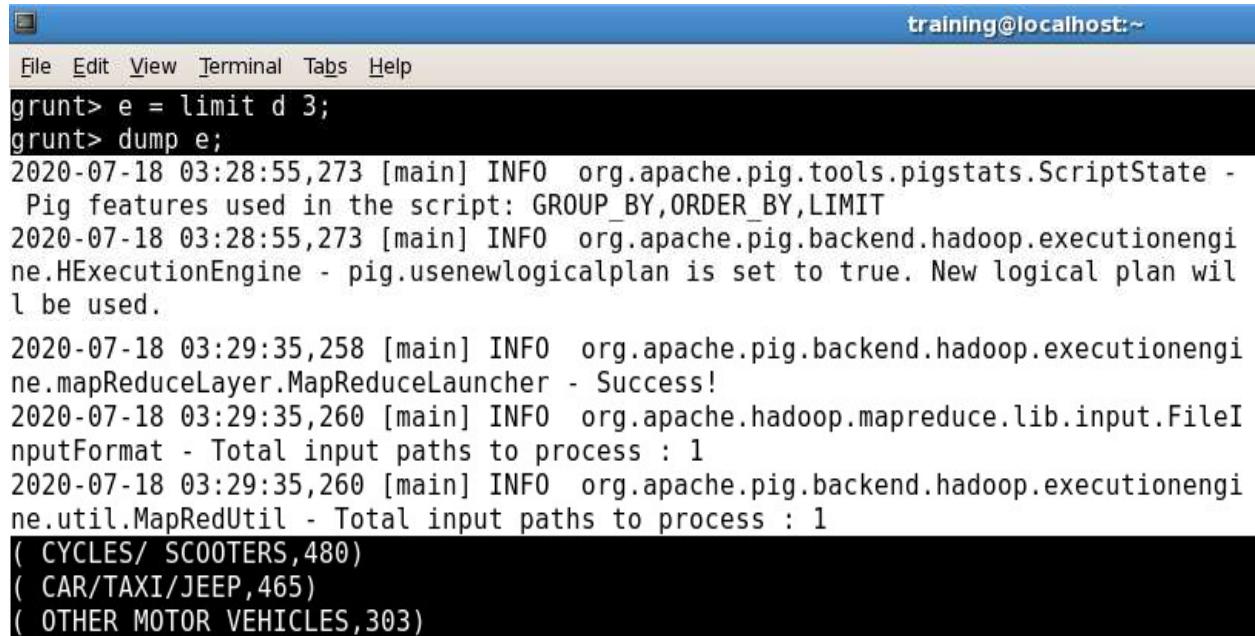
A screenshot of a terminal window titled "training@localhost:~". The window shows the following command-line session:

```
File Edit View Terminal Tabs Help
grunt> d = order c by c desc;
grunt> dump d;
2020-07-18 03:23:42,076 [main] INFO org.apache.pig.tools.pigstats.ScriptState -
Pig features used in the script: GROUP_BY,ORDER_BY
2020-07-18 03:23:42,076 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - pig.usenewlogicalplan is set to true. New logical plan will be used.
2020-07-18 03:23:42,124 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - (Name: d: Store(file:/tmp/temp216009587/tmp1195855272:org.
```

```
inputFormat - Total input paths to process : 1
2020-07-18 03:24:29,553 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
( CYCLES/ SCOOTERS,480)
( CAR/TAXI/JEEP,465)
( OTHER MOTOR VEHICLES,303)
( GOODS CARRYING VEHICLES (TRUCKS/TEMPO ETC),292)
( BUSES,254)
(TRUCKS/TEMPO ,124)
(OTHER VEHICLES,122)
(BUSES,103)
```

Step vi. Apply LIMIT operation as limit 3.

```
grunt> e = limit d 3;
grunt> dump e;
```



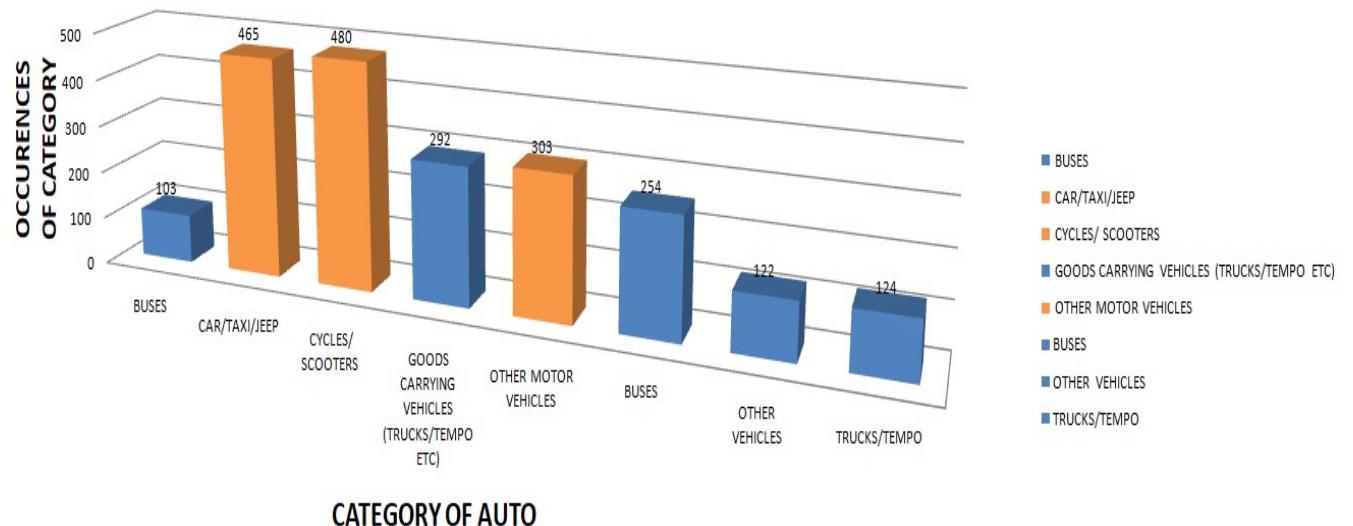
A screenshot of a terminal window titled "training@localhost:~". The window shows a sequence of Pig Latin commands and their execution logs. The commands include defining a relation 'e' with a LIMIT of 3, and then dumping it. The logs provide details about the execution environment, including the use of a new logical plan and the success of the map-reduce layer.

```
File Edit View Terminal Tabs Help
training@localhost:~
grunt> e = limit d 3;
grunt> dump e;
2020-07-18 03:28:55,273 [main] INFO org.apache.pig.tools.pigstats.ScriptState -
Pig features used in the script: GROUP_BY,ORDER_BY,LIMIT
2020-07-18 03:28:55,273 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - pig.usenewlogicalplan is set to true. New logical plan will be used.
2020-07-18 03:29:35,258 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-07-18 03:29:35,260 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-07-18 03:29:35,260 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
( CYCLES/ SCOOTERS,480)
( CAR/TAXI/JEEP,465)
( OTHER MOTOR VEHICLES,303)
```

VISUALIZATION

This column graph shows the graph plotted between CATEGORY OF AUTO and their OCCURENCES. The **THREE BAR** with **ORANGE COLOR** shows the top 3 category in the graph.

TOP 3 CATEGORY OF CHART GRAPH



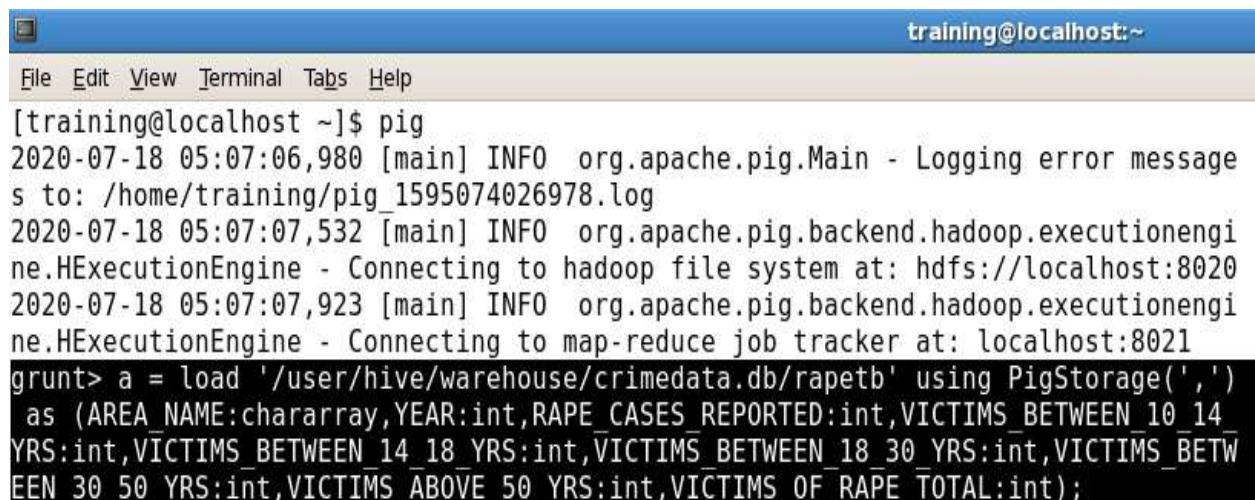
Query 9:- Year wise Total Rape case reported v/s Total rape victims.

CODING

Step i. Open Pig in MapReduce mode

Step ii. Create a bag with name ‘a’ and load the data present on Hadoop (dataset) into bag ‘a’

```
[training@localhost ~]$ pig
grunt> a = load
'/user/hive/warehouse/crimeData.db/rapetb' using
PigStorage(',') as
(AREA_NAME:chararray, YEAR:int, RAPE_CASES_REPORTED:int,
VICTIMS_BETWEEN_10_14_YRS:int, VICTIMS_BETWEEN_14_18_YRS:int,
VICTIMS_BETWEEN_18_30_YRS:int, VICTIMS_BETWEEN_30_50_YRS:int,
VICTIMS_ABOVE_50_YRS:int, VICTIMS_OF_RAPE_TOTAL:int);
```



The screenshot shows a terminal window with a blue header bar containing the text "training@localhost:~". Below the header is a menu bar with "File Edit View Terminal Tabs Help". The main area of the terminal contains the following text:

```
[training@localhost ~]$ pig
2020-07-18 05:07:06,980 [main] INFO org.apache.pig.Main - Logging error messages to: /home/training/pig_1595074026978.log
2020-07-18 05:07:07,532 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to hadoop file system at: hdfs://localhost:8020
2020-07-18 05:07:07,923 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine - Connecting to map-reduce job tracker at: localhost:8021
grunt> a = load '/user/hive/warehouse/crimeData.db/rapetb' using PigStorage(',') as (AREA_NAME:chararray, YEAR:int, RAPE_CASES_REPORTED:int, VICTIMS_BETWEEN_10_14_YRS:int, VICTIMS_BETWEEN_14_18_YRS:int, VICTIMS_BETWEEN_18_30_YRS:int, VICTIMS_BETWEEN_30_50_YRS:int, VICTIMS_ABOVE_50_YRS:int, VICTIMS_OF_RAPE_TOTAL:int);
```

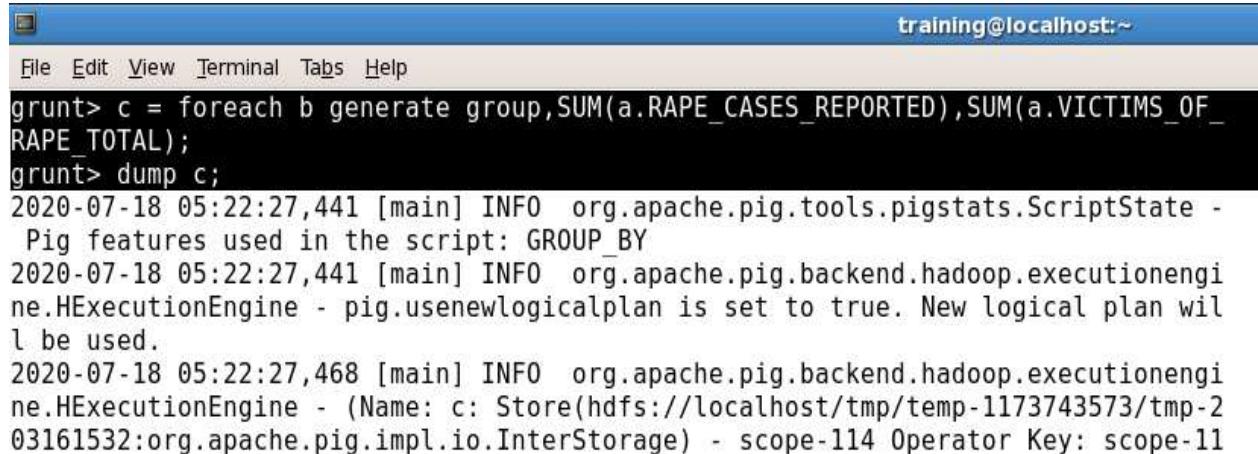
Step iii. Now perform GROUP operation by referencing with a column YEAR in a bag ‘b’

```
grunt> b = group a by YEAR;  
grunt> dump b;
```

```
grunt> b = group a by YEAR;  
grunt> dump b;  
,(TAMIL NADU,2014,686,50,139,411,66,0,666),(TAMIL NADU,2014,0,0,0,0,0,0),(TAMI  
L NADU,2014,686,50,139,411,66,0,666),(TRIPURA,2014,238,27,61,93,38,0,219),(TRIPU  
RA,2014,0,0,0,0,0,0),(TRIPURA,2014,238,27,61,93,38,0,219),(UTTAR PRADESH,2014,  
1563,162,246,862,250,0,1520),(UTTAR PRADESH,2014,4,1,2,1,0,0,4),(UTTAR PRADESH,2  
014,1559,161,244,861,250,0,1516),(UTTARAKHAND,2014,121,3,5,92,19,0,119),(UTTARAK  
HAND,2014,11,2,2,5,1,0,10),(UTTARAKHAND,2014,110,1,3,87,18,0,109),(WEST BENGAL,2  
014,2311,22,51,1779,459,0,2311),(WEST BENGAL,2014,4,1,0,2,1,0,4),(WEST BENGAL,20  
14,2307,21,51,1777,458,0,2307)})
```

Step iv. Perform SUM operation on ‘a’ in bag ‘c’

```
grunt> c = foreach b generate  
group,SUM(a.RAPE_CASES_REPORTED),SUM(a.VIC  
TIMS_OF_RAPE_TOTAL);  
grunt> dump c;
```



The screenshot shows a terminal window titled "training@localhost:~". The window contains the following text:

```
File Edit View Terminal Tabs Help  
training@localhost:~  
grunt> c = foreach b generate group,SUM(a.RAPE_CASES_REPORTED),SUM(a.VICTIMS_OF_  
RAPE_TOTAL);  
grunt> dump c;  
2020-07-18 05:22:27,441 [main] INFO org.apache.pig.tools.pigstats.ScriptState -  
Pig features used in the script: GROUP_BY  
2020-07-18 05:22:27,441 [main] INFO org.apache.pig.backend.hadoop.executionengi  
ne.HExecutionEngine - pig.usenewlogicalplan is set to true. New logical plan wil  
l be used.  
2020-07-18 05:22:27,468 [main] INFO org.apache.pig.backend.hadoop.executionengi  
ne.HExecutionEngine - (Name: c: Store(hdfs://localhost/tmp/temp-1173743573/tmp-2  
03161532:org.apache.pig.impl.io.InterStorage) - scope-114 Operator Key: scope-11
```

```

inputFormat - Total input paths to process : 1
2020-07-18 05:22:50,740 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1

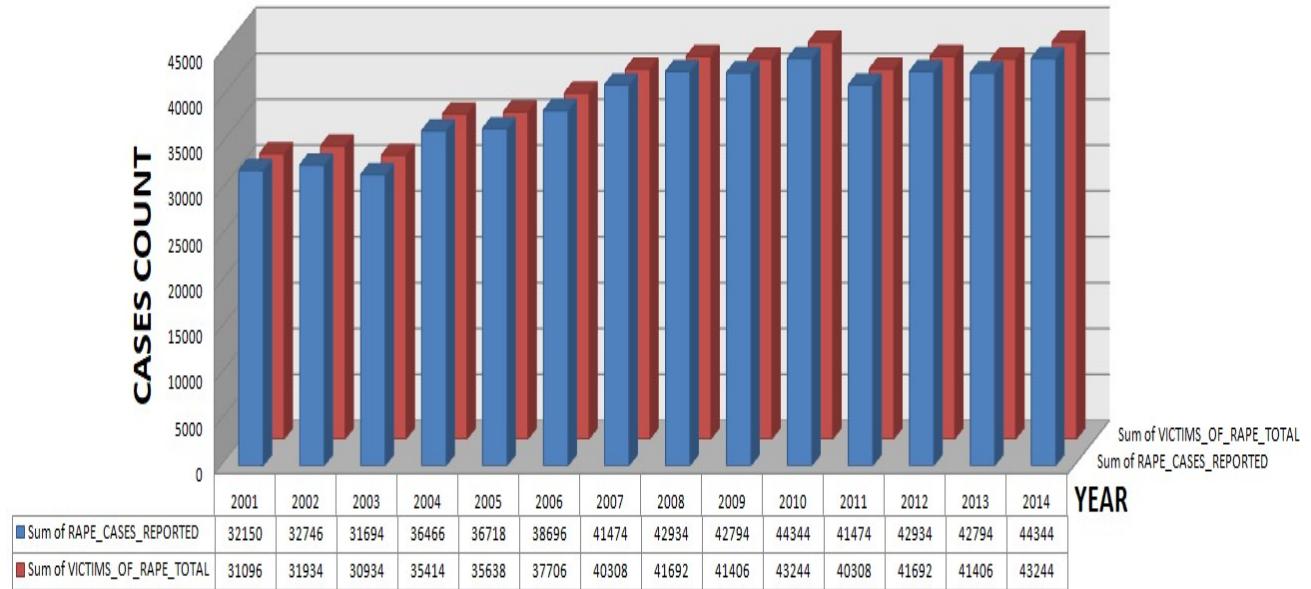
```

(2001,32150,31096)
(2002,32746,31934)
(2003,31694,30934)
(2004,36466,35414)
(2005,36718,35638)
(2006,38696,37706)
(2007,41474,40308)
(2008,42934,41692)
(2009,42794,41406)
(2010,44344,43244)
(2011,41474,40308)
(2012,42934,41692)
(2013,42794,41406)
(2014,44344,43244)

VISUALIZATION

This column graph shows the graphical comparison between RAPE CASES REPORTED and VICTIMS OF RAPE CASES.

RAPE CASES REPORTED V/S VICTIM CHART GRAPH



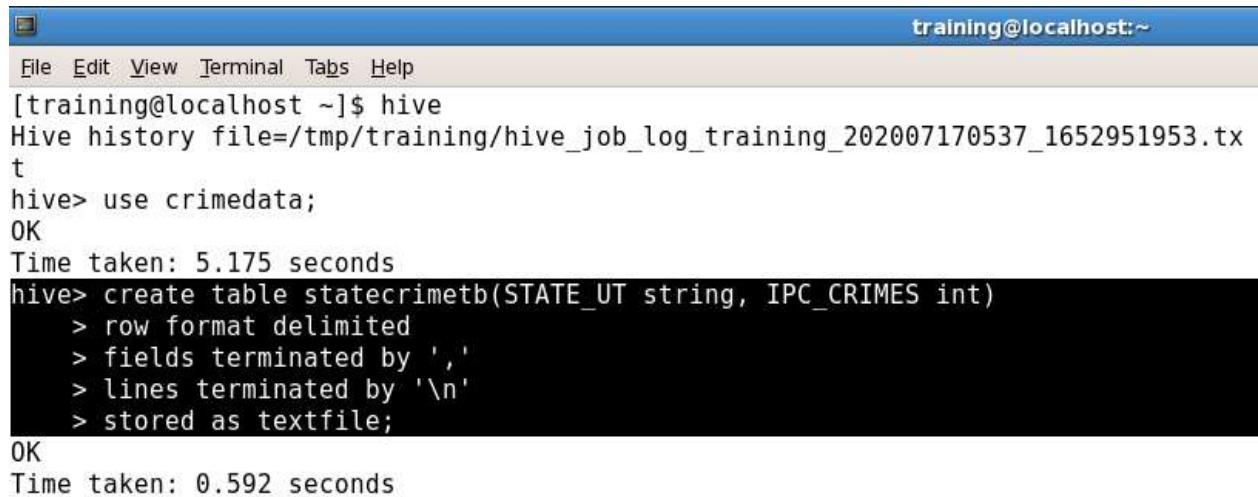
Query 10:- State wise percentage of IPC crimes recorded.

CODING

Step i. Open Hive

Step ii. Create a new table with name statecrimetb into previously created database i.e. crimedata

```
[training@localhost ~]$ hive
hive> use crimedata;
hive> create table statecrimetb(STATE_UT string,
IPC_CRIMES int)
> row format delimited
> fields terminated by ','
> lines terminated by '\n'
> stored as textfile;
```



The screenshot shows a terminal window titled "training@localhost:~". The window contains the following command-line session:

```
File Edit View Terminal Tabs Help
[training@localhost ~]$ hive
Hive history file=/tmp/training/hive_job_log_training_202007170537_1652951953.txt
hive> use crimedata;
OK
Time taken: 5.175 seconds
hive> create table statecrimetb(STATE_UT string, IPC_CRIMES int)
> row format delimited
> fields terminated by ','
> lines terminated by '\n'
> stored as textfile;
OK
Time taken: 0.592 seconds
```

Step iii. Insert data into statecrimetb from previously created crimetable

```
hive> insert overwrite table statecrimetb select  
STATE_UT,sum(TOTAL_IPC_CRIMES) from  
crimetable group by STATE_UT;  
hive> select * from statecrimetb;
```

```
hive> insert overwrite table statecrimetb select STATE_UT,sum(TOTAL_IPC_CRIMES)  
from crimetable group by STATE_UT;  
Total MapReduce jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks not specified. Estimated from input data size: 1  
In order to change the average load for a reducer (in bytes):  
    set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
    set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
    set mapred.reduce.tasks=<number>  
Starting Job = job_202007161753_0012, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007161753_0012  
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007161753_0012  
2020-07-17 05:40:05,807 Stage-1 map = 0%,  reduce = 0%  
2020-07-17 05:40:07,830 Stage-1 map = 100%,  reduce = 0%  
2020-07-17 05:40:18,978 Stage-1 map = 100%,  reduce = 33%  
hive> select * from statecrimetb;  
OK  
A & N ISLANDS      18230  
A&N ISLANDS        2288  
ANDHRA PRADESH     4614898  
ARUNACHAL PRADESH   67222  
ASSAM              1606674  
BIHAR              3413366  
CHANDIGARH         96520  
CHHATTISGARH       1340494  
D & N HAVELI        9476  
D&N HAVELI         1106  
DAMAN & DIU          6808  
DELHI UT            1742776  
GOA                80142  
GUJARAT            3316696  
HARYANA            1472820  
HIMACHAL PRADESH    363354  
JAMMU & KASHMIR      616670  
JHARKHAND           1033274
```

Step iv. Perform SUM operation over the column
IPC_CRIMES of table statecrimetb

```
hive> select sum(IPC_CRIMES) from statecrimetb;
```

```
hive> select sum(IPC_CRIMES) from statecrimetb;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202007161753_0020, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007161753_0020
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007161753_0020
2020-07-17 06:07:27,350 Stage-1 map = 0%,  reduce = 0%
2020-07-17 06:07:29,614 Stage-1 map = 100%,  reduce = 0%
2020-07-17 06:07:39,276 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_202007161753_0020
OK
58206005
Time taken: 17.112 seconds
```

Step v. Apply PERCENTAGE over the column
IPC_CRIMES of statecrimetb

```
hive> select STATE_UT,IPC_CRIMES*100/58206005
from statecrimetb;
```

training@localhost:~

File Edit View Terminal Tabs Help

```
hive> select STATE UT,IPC CRIMES*100/58206005 from statecrimetb;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_202007161753_0019, Tracking URL = http://localhost:50030/jobdetails.jsp?jobid=job_202007161753_0019
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -kill job_202007161753_0019
2020-07-17 06:06:48,705 Stage-1 map = 0%,  reduce = 0%
2020-07-17 06:06:50,880 Stage-1 map = 100%,  reduce = 0%
2020-07-17 06:06:51,886 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_202007161753_0019
OK
A & N ISLANDS    0.03131979251968933
A&N ISLANDS      0.003930865896053165
ANDHRA PRADESH   7.928559948410821
ARUNACHAL PRADESH 0.11548980212608648
ASSAM     2.760323440854599
BIHAR     5.864284965099391
CHANDIGARH        0.16582481481077424
CHHATTISGARH       2.303016673279673
D & N HAVELI      0.016280107181381027
D&N HAVELI        0.0019001475878648603
DAMAN & DIU        0.011696387683710641
DELHI UT          2.9941515484527756
GOA        0.137686824581072
GUJARAT      5.698202444919557
```

HARYANA 2.5303574777207265
HIMACHAL PRADESH 0.6242551777947997
JAMMU & KASHMIR 1.0594611329191894
JHARKHAND 1.7752017167300864
KARNATAKA 5.920468171625934
KERALA 5.852282767044397
LAKSHADWEEP 0.0029481494220398737
MADHYA PRADESH 9.874149583019827
MAHARASHTRA 9.345070152126057
MANIPUR 0.14130157189107895
MEGHALAYA 0.1098477725794787
MIZORAM 0.10253237616977835
NAGALAND 0.05254097064383649
ODISHA 2.733810712485765
PUDUCHERRY 0.19217432977920404
PUNJAB 1.5766105232613026
RAJASTHAN 7.789436845906192
SIKKIM 0.030069062461854236
TAMIL NADU 8.268875350575941
TELANGANA 0.31029444470549045
TRIPURA 0.2226505667241722
UTTAR PRADESH 7.983990655259711
UTTARAKHAND 0.4188021493658601
WEST BENGAL 5.050200576383828

Time taken: 9.616 seconds

VISUALIZATION

This Pie-chart graph shows percentage wise representation of IPC CRIMES in STATE/UT.

STATE/UT - CRIMES PERCENTAGE CHART GRAPGH

